

Convergence of Random Variables

Here, in this part, we will consider infinite sequences of random variables and their limiting (asymptotic) behavior. In fact, one of the subfield of the study in statistics is the Large Sample Statistics, or the Asymptotic Statistics, where the asymptotic behavior of statistics and tests are studied. The idea is that when one wants to consider a large sample, say, we have 10^{10} observations, then it is sometimes much easier not to consider what happen when $n = 10^{10}$, but to consider the limit when $n \rightarrow +\infty$. Say, if we will ask every person on Earth to choose an integer number from $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$, and denote the choice of n -th person by X_n , then the average number chosen will be

$$\bar{X}_N = \frac{X_1 + X_2 + \dots + X_N}{N},$$

where $N \approx 7.608\text{bln}$ as of 2018 March, the total population on the Earth¹. Now, instead of calculating this fraction, we can calculate the limit of that: because of the LLN (see forthcoming), we will have that

$$\bar{X}_n \rightarrow \mathbb{E}(X_1) = \frac{1}{10} \cdot 0 + \frac{1}{10} \cdot 1 + \frac{1}{10} \cdot 2 + \dots + \frac{1}{10} \cdot 9 = \frac{45}{10} = 4.5,$$

assuming that the choice of digits is uniform. Then we can state that $\bar{X}_N \approx 4.5$, since N is large enough².

In parametric statistics, one is considering a sample of size n , and estimating an unknown quantity (parameter) θ through a r.v. variable θ_n , which is an estimate for that unknown θ . Here θ_n depends on the sample size n (we will consider point estimation in the next part). Usually, we want to study the asymptotic behaviour of θ_n , the behaviour of θ_n for large samples: for "good" estimates we want to have that θ_n approaches θ in some sense.

Here, in this part we will define different types of convergences and investigate their properties and relationship.

We will assume that we are given a sequence of r.v. $X_1, X_2, \dots, X_n, \dots$ (infinite number of r.v.). Usually, we will assume that all these r.v. are defined on the same Experiment, on the same Probability Space. That is, mostly we will assume that $X_n = X_n(\omega)$, where $\omega \in \Omega$ for some fixed Ω . Only when defining the Convergence in Distributions notion, we will rest(skip?) this assumption.

As examples of sequences of r.v. we can consider:

EXAMPLE, SEQUENCES OF R.V.:

- Let X_n be the daily closing price for one Facebook Inc. stock at the day n calculated from now. Then X_n will be a sequence of r.v. . What is the Sample Space behind each X_k ? We can assume that the sample space is the set of all possible market scenarios for the rest of the world.

¹Well, of course, newborn babies will not choose for your experiment numbers ☺ But let's assume, for the effect of the presentation.

Btw, can you estimate how many people have ever lived on Earth?

²Well, mathematically, this is not correct, but in practice ... ☺

Here one fact is notable: X_n -s will not be independent, the price tomorrow will depend somehow from the today's price. Here, in our course, we will mainly consider IID sequences, i.e. sequences X_k such that all X_k -s are Identically Distributed (have the same distribution) and are Independent.

- Assume we are tossing a coin infinitely many times³, $X_n = 0$, if the n -th toss resulted in tails, and $X_n = 1$, if heads.
- Let X_n be the number of insurance claims for some insurance company for the day n calculated from today, and let Y_n be the claim size for that day. Then X_n and Y_n are sequences of r.v.

A naive way to think about the sequence of r.v.'s on the same probability space - say, we have (countably) infinite number of persons numbered by $1, 2, 3, \dots$, watching after the result of the experiment. Everybody has his/her own function to calculate for any possible outcome (say, the Sample Space is $\Omega = [1, 6]$, the first person will calculate the square of ω , i.e. $X_1(\omega) = \omega^2$, the second person will calculate the inverse of ω , $X_2(\omega) = \frac{1}{\omega}$ etc.). The only unsure thing is - which outcome will happen. So before doing the experiment, the choice of every person is a r.v..

8.1 Convergence of a sequence of r.v.'s

Assume now we have a sequence of r.v. X_n , $n \in \mathbb{N}$, defined on the same Probability Space $(\Omega, \mathcal{F}, \mathbb{P})$. We want to define some convergence notions for that sequence.

Assume also X is a r.v. on the same space $(\Omega, \mathcal{F}, \mathbb{P})$.

Definition 8.1. We will say that $X_n \rightarrow X$ almost sure, and we will write $X_n \rightarrow X$ a.s. or $X_n \xrightarrow{\text{a.s.}} X$, if

$$\mathbb{P}\left(\omega \in \Omega : \lim_{n \rightarrow +\infty} X_n(\omega) = X(\omega)\right) = 1,$$

or, for short,

$$\mathbb{P}(X_n \rightarrow X) = 1$$

Equivalently, we can write

$$X_n \xrightarrow{\text{a.s.}} X \quad \text{iff} \quad \mathbb{P}(X_n \not\rightarrow X) = 0.$$

You can think like this: X_n is a r.v., so it is a function of ω . If we fix an ω , then $X_n(\omega)$ becomes a numerical sequence. For this numerical sequence, we can consider its limit. If the limit exists at ω and is equal to $X(\omega)$, then ω is from the set $X_n \rightarrow X$. Now, the a.s. convergence idea is that the set (of all ω -s) where X_n does not tend to X is of probability 0.

Definition 8.2. We will say that $X_n \rightarrow X$ in Probability, and we will write $X_n \xrightarrow{\mathbb{P}} X$, if

$$\text{for any } \varepsilon > 0, \quad \mathbb{P}(|X_n - X| \geq \varepsilon) \rightarrow 0, \quad \text{when } n \rightarrow \infty.$$

Equivalently, we can write

$$X_n \xrightarrow{\mathbb{P}} X \quad \text{iff} \quad \mathbb{P}(|X_n - X| < \varepsilon) \rightarrow 1 \text{ for any } \varepsilon > 0.$$

Sometimes we will use the first form, with $\mathbb{P}(|X_n - X| \geq \varepsilon) \rightarrow 0$, and sometimes we will use the equivalent form $\mathbb{P}(|X_n - X| < \varepsilon) \rightarrow 1$.

Another notion of convergence is the Quadratic Mean convergence:

Definition 8.3. We will say that $X_n \rightarrow X$ in Quadratic Mean or in L^2 (or in Mean Square Sense), and we will write $X_n \xrightarrow{L^2} X$ or $X_n \xrightarrow{qm} X$, if

$$\mathbb{E}((X_n - X)^2) \rightarrow 0, \quad \text{when } n \rightarrow \infty.$$

In the above definitions, it is important that all random variables are defined on the same probability space. This is necessary to calculate $X_n(\omega)$ and $X(\omega)$ for the same ω .

On the other hand, for the CLT, this is too restrictive, since CLT works for a broad range of r.v.s, not necessarily defined on the same probability space. So we give another type of convergence notion, which will work also in that case. The idea is even if the random variables are defined on completely different spaces, their CDF's are just real valued functions defined on \mathbb{R} , so we can talk about the convergence of their CDF's.

So now we assume that X_n and X are arbitrary r.v.'s, not necessarily defined on the same probability space, and $F_{X_n}(x)$ and $F_X(x)$ are their CDF's, respectively.

Definition 8.4. We will say that $X_n \rightarrow X$ in Distribution, and we will write $X_n \xrightarrow{D} X$, if

$$F_{X_n}(x) \rightarrow F_X(x) \quad \text{when } n \rightarrow \infty \text{ at any point of continuity } x \text{ of } F_X(x).$$

REMARK, CONVERGENCE IN DISTRIBUTIONS: The above definition says that $X_n \xrightarrow{D} X$ iff

$$\mathbb{P}(X_n \leq x) \rightarrow \mathbb{P}(X \leq x),$$

at any point of continuity of $\mathbb{P}(X \leq x)$. So convergence in distribution is the convergence of corresponding probabilities.

In general, if $X_n \xrightarrow{D} X$, then for many subsets $A \subset \mathbb{R}$, one will have

$$\mathbb{P}(X_n \in A) \rightarrow \mathbb{P}(X \in A).$$

In fact, the convergence will hold for all **continuity sets** A of random variable X , see https://en.wikipedia.org/wiki/Continuity_set. And, in particular, for any $a, b \in \mathbb{R}$ with $\mathbb{P}(X = a) = \mathbb{P}(X = b) = 0$, we will have

$$\mathbb{P}(a \leq X_n \leq b) \rightarrow \mathbb{P}(a \leq X \leq b), \quad n \rightarrow +\infty.$$

Proposition 8.1. Assume X_n is a sequence of r.v. and X is a r.v. . Then

$$X_n \xrightarrow{D} X \quad \text{iff} \quad \mathbb{E}(h(X_n)) \rightarrow \mathbb{E}(h(X)) \quad \text{for any continuous and bounded function } h.$$

The following chart represents the relationship between different types of convergences:

Convergence in QM \Rightarrow Convergence in Probability \Rightarrow Convergence in Distribution
 Convergence AS

REMARK, RELATIONSHIP BETWEEN DIFFERENT TYPES OF CONVERGENCES: It can be proved, that the inverse implications in the above diagram are not true, in general.

!!! Give examples!

But, we can state some partial results:

- If $X_n \xrightarrow{D} X$ and $X \equiv \text{constant}$, then $X_n \xrightarrow{P} X$;
- If $X_n \xrightarrow{P} X$, then there is a subsequence of natural numbers n_k such that $X_{n_k} \xrightarrow{a.s.} X$;
- If $X_n \xrightarrow{P} X$ and there exists a number M such that $\mathbb{P}(|X_n| \leq M) = 1$ for any n , then $X_n \xrightarrow{q.m.} X$.

EXAMPLE, CONVERGENCE IN PROBABILITY BUT NOT A.S.: This is a classical example of a sequence of r.v. X_n such that X_n tends to 0 in probability but not a.s. .

We consider the Sample Space $\Omega = [0, 1]$ with a probability $\mathbb{P}([a, b]) = b - a$, for $[a, b] \subset \Omega$. We define the sequence X_n in the following way:

$$\begin{aligned} X_1(\omega) &\equiv 1, & X_2(\omega) &= \begin{cases} 1, & \omega \in [0, \frac{1}{2}] \\ 0, & \text{otherwise} \end{cases} & X_3(\omega) &= \begin{cases} 0, & \omega \in [0, \frac{1}{2}] \\ 1, & \text{otherwise} \end{cases} \\ X_4(\omega) &= \begin{cases} 1, & \omega \in [0, \frac{1}{2^2}] \\ 0, & \text{otherwise} \end{cases} & X_5(\omega) &= \begin{cases} 1, & \omega \in [\frac{1}{2^2}, \frac{2}{2^2}] \\ 0, & \text{otherwise} \end{cases} \\ X_6(\omega) &= \begin{cases} 1, & \omega \in [\frac{2}{2^2}, \frac{3}{2^2}] \\ 0, & \text{otherwise} \end{cases} & X_7(\omega) &= \begin{cases} 1, & \omega \in [\frac{3}{2^2}, 1] \\ 0, & \text{otherwise} \end{cases} \end{aligned}$$

and so on. In fact, to construct r.v. X_n for $n = 8, 9, 10, \dots, 15$, we divide $[0, 1]$ into 8 equal-length intervals and define

$$X_{2^3+k}(\omega) = \begin{cases} 1, & \omega \in [\frac{k}{2^3}, \frac{k+1}{2^3}] \\ 0, & \text{otherwise} \end{cases} \quad \text{for all } k = 0, 1, 2, \dots, 7.$$

And, in general, for any $m = 0, 1, 2, \dots$ and $k = 0, 1, 2, \dots, 2^m - 1$

$$X_{2^m+k}(\omega) = \begin{cases} 1, & \omega \in [\frac{k}{2^m}, \frac{k+1}{2^m}] \\ 0, & \text{otherwise} \end{cases}$$

The graphs of the first few function (r.v.-s) X_k are give in Fig.

Now, first we prove that at any $\omega \in [0, 1]$, $X_n(\omega)$ is divergent. To prove this, it is enough to note that the numerical sequence $X_n(\omega)$ consists of infinitely many 0-s and infinitely many 1-s, so it cannot converge. This, particularly, means that

$$\mathbb{P}(\omega \in \Omega | X_n(\omega) \text{ converges}) = 0,$$

that is, X_n diverges a.s. (in fact, diverges surely, at any point).

On the other hand, we can see that X_n converges to 0 in the sense of Probabilities. To see this, we fix a positive $\varepsilon < 1$, and calculate $\mathbb{P}(|X_n - 0| \geq \varepsilon) \stackrel{X_n \text{ is non-negative}}{=} \mathbb{P}(X_n \geq \varepsilon)$. Clearly,

$$\mathbb{P}(X_1 \geq \varepsilon) = \mathbb{P}([0, 1]) = 1, \quad \mathbb{P}(X_2 \geq \varepsilon) = \mathbb{P}([0, \frac{1}{2}]) = \frac{1}{2}, \quad \mathbb{P}(X_3 \geq \varepsilon) = \mathbb{P}([\frac{1}{2}, 1]) = \frac{1}{2},$$

$$\mathbb{P}(X_4 \geq \varepsilon) = \mathbb{P}([0, \frac{1}{2^2}]) = \frac{1}{2^2}, \mathbb{P}(X_5 \geq \varepsilon) = \mathbb{P}([\frac{1}{2^2}, \frac{2}{2^2}]) = \frac{1}{2^2}, \dots$$

and so on, so the sequence $\mathbb{P}(X_n \geq \varepsilon)$ has the form

$$1, \frac{1}{2}, \frac{1}{2}, \frac{1}{2^2}, \frac{1}{2^2}, \frac{1}{2^2}, \frac{1}{2^2}, \frac{1}{2^3}, \frac{1}{2^3}, \dots$$

and this clearly tends to 0. This means that $X_n \xrightarrow{\mathbb{P}} 0$.

In fact, for convergence a.s., most of the algebraic operations with sequences remain true. But this is **not true for the convergence in Distributions**.

Proposition 8.2. (Uniqueness of a limit)

- Assume $X_n \xrightarrow{\text{a.s.}} X$ and $X_n \xrightarrow{\text{a.s.}} Y$. Then $X = Y$ a.s., i.e., $\mathbb{P}(X = Y) = 1$ (or, in other words, $\mathbb{P}(X \neq Y) = 0$).
- Assume $X_n \xrightarrow{\mathbb{P}} X$ and $X_n \xrightarrow{\mathbb{P}} Y$. Then $X = Y$ a.s..
- Assume $X_n \xrightarrow{\text{q.m.}} X$ and $X_n \xrightarrow{\text{q.m.}} Y$. Then $X = Y$ a.s..

REMARK, NON-UNIQUENESS OF THE LIMIT IN THE CONVERGENCE IN DISTRIBUTIONS CASE: It is important that the above proposition is not true for the convergence in Distributions, that is, it can happen that $X_n \xrightarrow{D} X$ and $X_n \xrightarrow{D} Y$, where $X \neq Y$ at most of the points, and even, $X \neq Y$ a.s.

For example, consider some r.v. with symmetric distribution around 0, e.g., $X \sim \mathcal{N}(0, 1)$. Take $X_n = X$ for all n . Clearly, $X_n \xrightarrow{D} X$. But, also $X_n \xrightarrow{D} -X$. To see this, we need to prove that for any $x \in \mathbb{R}$

$$F_{X_n}(x) \rightarrow F_{-X}(x)$$

(as $-X$ is a continuous r.v., so F_{-X} is continuous everywhere). But it is easy to prove that $F_{-X}(x) = F_X(x)$ for any x , because of the symmetry of the distribution of X :

$$F_{-X}(x) = \mathbb{P}(-X \leq x) = \mathbb{P}(X \geq -x) \stackrel{\text{because of the symmetry}}{=} \mathbb{P}(X \leq x) = F_X(x).$$

Now, $X_n \xrightarrow{D} X$ and $X_n \xrightarrow{D} -X$. Maybe $X = -X$ a.s.? Of course, no! Because

$$\mathbb{P}(X = -X) = \mathbb{P}(2X = 0) = \mathbb{P}(X = 0) = 0,$$

so $X \neq -X$ a.s..

So we can have that the same sequence of r.v. X_n can have several (in fact, a lot of) different limits in the sense of Distributions convergence, but, important point is that **the distribution of all limiting r.v.'s will be the same**. Say, in our above case, X and $-X$ both have $\mathcal{N}(0, 1)$ distribution.

Proposition 8.3. Assume $X_n \xrightarrow{\text{a.s.}(P)} X$ and $Y_n \xrightarrow{\text{a.s.}(P)} Y$. Then

- a. $X_n + Y_n \xrightarrow{\text{a.s.}(P)} X + Y;$
- b. $X_n \cdot Y_n \xrightarrow{\text{a.s.}(P)} X \cdot Y;$
- c. If $g \in C(\mathbb{R})$, then $g(X_n) \xrightarrow{\text{a.s.}(P)} g(X)$

Proposition 8.4. Assume $X_n \xrightarrow{L^2} X$ and $Y_n \xrightarrow{L^2} Y$. Then $X_n + Y_n \xrightarrow{L^2} X + Y$.

The same properties are not true, in general, for the convergence in Distributions. For example, in the general case, if $X_n \xrightarrow{D} X$ and $Y_n \xrightarrow{D} Y$, then not necessarily⁴ $X_n + Y_n \xrightarrow{D} X + Y$.

Example: Assume $Z \sim N(0, 1)$, $X_n = -Z + \frac{1}{n}$, $Y_n = -X_n$. Then $X_n \xrightarrow{D} Z$ and $Y_n \xrightarrow{D} Z$, but $X_n + Y_n = 0$. ■

The next Proposition gives some properties we can use when dealing with the convergence in distributions.

Theorem 8.1 (Slutsky's Theorem). Assume $X_n \xrightarrow{D} X$ and $Y_n \xrightarrow{P} c$, where $c \in \mathbb{R}$ is a constant. Then

- a. $X_n + Y_n \xrightarrow{D} X + c;$
- b. $X_n \cdot Y_n \xrightarrow{D} c \cdot X.$

Theorem 8.2 (Continuous Mapping Theorem). Assume $X_n \xrightarrow{D} X$ and $g \in C(\mathbb{R})$. Then $g(X_n) \xrightarrow{D} g(X)$.

Sometimes, when using the convergence in distributions, we will write something like $X_n \xrightarrow{D} N(\mu, \sigma^2)$ instead of writing $X_n \xrightarrow{D} X$ for some $X \sim N(\mu, \sigma^2)$. This is because X is not important, its distribution is important in case of the convergence in Distributions.

8.2 Examples of convergence

Usually, one will have 2 types of examples of sequence of r.v., and we will give that 2 types of examples, explaining or checking their convergence.

The first type of examples are when we describe r.v. X_n **explicitly**, i.e., as a function of ω . So in that case we talk about a **concrete example of a sequence**. Usually, the explicit form of X_n will not be available: say, if X_n will be the claim size for the day n , then nobody will give a formula like $X_n(\omega) = 100n + \omega^2$ - even it is hard to describe what is ω here: usually, ω represents the scenario, and it will be non-numerical, and will be undescrivable (??). So usually, this type of examples are introduced to explain the ideas of convergence notions.

The second type of examples concern r.v. X_n that are given by distribution. So we talk about the distribution of X_n , **without explicitly describing** X_n . Say, we are talking about $X_n \sim \text{Bernoulli}(0.5)$: this means that for any fixed n , X_n is some function (defined on some Sample Space Ω , which is not shown anywhere), which takes the values 0 and 1 (and we are not specifying for which ω the value is 0, and for which - 1), with probabilities 0.5 both (so for "half" of the values⁵ ω , the value $X_n(\omega)$ will be 0, and for all others will be 1). Of course, we can have different sets for taking the values 0 and 1 for different n .

Or, say, we will consider $X_n \sim \text{Unif}([0, \frac{1}{n}])$. In this case we will calculate the limit of X_n in one of the senses above **irrespective** to the concrete realisation, concrete value of X_n . I am stressing

⁴The problem is that we need some information about the Joint Distribution of (X_n, Y_n) .

⁵I am using here "half" of the values meaning that the probability of the set of that ω -s is 0.5.

this because we can have a lot of r.v.'s X_n having the same distribution, so $X_n \sim \text{Unif}([0, \frac{1}{n}])$ is not unique⁶! But the convergence will hold in any case!

Now, the examples.

8.2.1 R.V. described Explicitly

EXAMPLE, CONVERGENCE OF A R.V. SEQUENCE: Assume we are rolling a fair die, $\Omega = \{1, 2, 3, 4, 5, 6\}$ is the Sample Space, and assume X_n is calculating the result shown on the die divided by n , i.e.

$$X_n(\omega) = \frac{\omega}{n}, \quad \omega \in \Omega.$$

We can describe this as: if the result will be 3, then, say, the first person is calculating $X_1 = \frac{3}{1} = 3$, the second person calculates $X_2 = \frac{3}{2}$, for the third one $X_3 = \frac{3}{3} = 1$ etc.

Let us prove that $X_n \rightarrow 0$ in all four senses. Let us denote $X \equiv 0$, and prove that $X_n \rightarrow X$ is all senses.

Almost Sure Convergence: To prove that $X_n \xrightarrow{\text{a.s.}} X$, we need to see for which $\omega \in \Omega$ we will have

$$X_n(\omega) \rightarrow X(\omega).$$

Now, if $\omega \in \Omega$ is fixed, then

$$X_n(\omega) = \frac{\omega}{n} \rightarrow 0 = X(\omega).$$

This means that for any ω , $X_n(\omega) \rightarrow X(\omega)$. Then the set $\{\omega \in \Omega : X_n(\omega) \rightarrow X(\omega)\} = \Omega$, so

$$\mathbb{P}(X_n \rightarrow X) = \mathbb{P}(\Omega) = 1.$$

Quadratic Mean Convergence: To prove that $X_n \xrightarrow{\text{qm}} X$, we need to prove that

$$\mathbb{E}((X_n - X)^2) \rightarrow 0.$$

To this end, we need to calculate the expected value on the left:

$$\mathbb{E}((X_n - X)^2) = \mathbb{E}((X_n)^2).$$

Now, let us build the distribution (PMF) of X_n :

Values of X_n	$\frac{1}{n}$	$\frac{2}{n}$	$\frac{3}{n}$	$\frac{4}{n}$	$\frac{5}{n}$	$\frac{6}{n}$
$\mathbb{P}(X_n = x)$	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{1}{6}$

Then,

$$\mathbb{E}((X_n - X)^2) = \mathbb{E}((X_n)^2) = \frac{1}{6} \cdot \left(\frac{1}{n}\right)^2 + \frac{1}{6} \cdot \left(\frac{2}{n}\right)^2 + \frac{1}{6} \cdot \left(\frac{3}{n}\right)^2 + \frac{1}{6} \cdot \left(\frac{4}{n}\right)^2 + \frac{1}{6} \cdot \left(\frac{5}{n}\right)^2 + \frac{1}{6} \cdot \left(\frac{6}{n}\right)^2 \rightarrow 0,$$

since each term goes to 0 as $n \rightarrow \infty$.

⁶Give explicit different examples of $X \sim \text{Unif}[0, 1]$, i.e., construct one or many probability spaces Ω and different formulas for $X(\omega)$ such that $X \sim \text{Unif}[0, 1]$

Convergence in Probability: Well, in fact, we can use the property that convergence a.s. (or in Quadratic Mean) implies Convergence in Probability. But let us prove the Convergence in Probability using the very definition. So to prove that $X_n \xrightarrow{P} X$, we need to prove that

$$\text{for any } \varepsilon > 0, \quad \mathbb{P}(|X_n - X| \geq \varepsilon) \rightarrow 0, \quad \text{when } n \rightarrow \infty.$$

Let us fix an $\varepsilon > 0$. Then,

$$\mathbb{P}(|X_n - X| \geq \varepsilon) \stackrel{X \equiv 0, X_n \geq 0}{=} \mathbb{P}(X_n \geq \varepsilon) \stackrel{X_n(\omega) = \frac{\omega}{n}}{=} \mathbb{P}\left(\omega \in \Omega : \frac{\omega}{n} \geq \varepsilon\right).$$

Now if $n > \frac{6}{\varepsilon}$, then we will have $n > \frac{\omega}{\varepsilon}$ for any $\omega \in \Omega = \{1, 2, 3, 4, 5, 6\}$. So no ω will satisfy $\frac{\omega}{n} \geq \varepsilon$. This means that for $n > \frac{6}{\varepsilon}$,

$$\mathbb{P}\left(\omega \in \Omega : \frac{\omega}{n} \geq \varepsilon\right) = 0 \rightarrow 0.$$

Accordingly,

$$\mathbb{P}(|X_n - X| \geq \varepsilon) = \mathbb{P}\left(\omega \in \Omega : \frac{\omega}{n} \geq \varepsilon\right) \rightarrow 0.$$

Convergence in Distributions: Again, we can use the property that Convergence in Probability implies Convergence in Distributions. But let us prove the Convergence in Distributions using the definition. To that end, we need to calculate the CDF's of X_n , $F_{X_n}(x)$ and of X , $F_X(x)$, and show that

$$F_{X_n}(x) \rightarrow F_X(x) \quad \text{when } n \rightarrow \infty \quad \text{at any point of continuity } x \text{ of } F_X(x).$$

First, we have $X \equiv 0$, so

$$F_X(x) = \mathbb{P}(X \leq x) = \begin{cases} 1, & x \geq 0 \\ 0, & x < 0. \end{cases}$$

Note that the only discontinuity point of F_X is $x = 0$. So we need to check the above (CDF's) convergence at any point, except $x = 0$ (maybe we will have convergence at that point too, but this is not necessary to have Convergence in Distributions: it is enough to check for continuity points of F_X).

Now, about the CDF of X_n . Hope you remember how to construct CDF's for discrete random variables, and you can easily calculate from the PMF of X_n that

$$F_{X_n}(x) = \mathbb{P}(X_n \leq x) = \begin{cases} 0, & x < \frac{1}{n} \\ \frac{1}{6}, & \frac{1}{n} \leq x < \frac{2}{n} \\ \frac{2}{6}, & \frac{2}{n} \leq x < \frac{3}{n} \\ \frac{3}{6}, & \frac{3}{n} \leq x < \frac{4}{n} \\ \frac{4}{6}, & \frac{4}{n} \leq x < \frac{5}{n} \\ \frac{5}{6}, & \frac{5}{n} \leq x < \frac{6}{n} \\ 1, & x \geq \frac{6}{n} \end{cases}$$

Now, to prove that $F_{X_n}(x) \rightarrow F_X(x)$, for any $x \neq 0$, let us consider cases. Fix $x \neq 0$. If $x < 0$, then $F_{X_n}(x) = F_X(x) = 0$, for any n , so obviously, $F_{X_n}(x) = 0 \rightarrow 0 = F_X(x)$. The other case is when $x > 0$, and in this case we can calculate the value of F_X : $F_X(x) = 1$. Since x is fixed, then if $n > \frac{6}{x}$, we will have $x > \frac{6}{n}$, so we will have $F_{X_n}(x) = 1$, if $n > \frac{6}{x}$. Hence, $F_{X_n}(x) \rightarrow 1 = F_X(x)$.

Visually, the graphs of F_X and F_{X_n} for some values of n are given in Fig. 8.1-8.4. Hope it is visible for you that F_{X_n} approaches F_X at every point $x \neq 0$.

Joy and Happiness!

Extra Joy and Happiness: For the complete happiness, try to answer the following question: What about the point $x = 0$? Is it true that $F_{X_n}(0) \rightarrow F_X(0)$?




Fig. 8.1: Graph of the CDF of X , $F_X(x)$






Fig. 8.2: Graph of the CDF of X_1 , $F_{X_1}(x)$

Fig. 8.3: Graph of the CDF of X_3 , $F_{X_3}(x)$ Fig. 8.4: Graph of the CDF of X_{20} , $F_{X_{20}}(x)$

EXAMPLE, CONVERGENCE IN DIFFERENT SENSES: Assume we are again rolling a fair die, $\Omega = \{1, 2, 3, 4, 5, 6\}$, and assume X_n is given by

$$X_n(\omega) = \omega + \frac{n}{n+1}, \quad \omega \in \Omega.$$

Then $X_n(\omega) \rightarrow X(\omega)$ in all four senses, where $X(\omega) = \omega + 1$.

Surely, you can solve this now!

EXAMPLE, CONVERGENCE IN DIFFERENT SENSES: We have considered the sequence $X_n(\omega) = \omega + \frac{1}{n}$, $\omega \in \Omega = [0, 1]$ (with the length as a probability) during our lecture. Consider also $X_n(\omega) = \frac{n \cdot \omega^2}{n+3} - \frac{2}{n}$.

EXAMPLE, CONVERGENCE IN DIFFERENT SENSES: Assume now our experiment is the following: we are choosing a point at random, uniformly, from $\Omega = [0, 1]$. For any $n \in \mathbb{N}$, we define X_n to be

$$X_n(\omega) = \left(\omega + \frac{1}{n} \right)^2, \quad \forall \omega \in [0, 1], \quad \forall n \in \mathbb{N}.$$

Now, intuitively, $X_n(\omega) \rightarrow X(\omega)$, where $X(\omega) = \omega^2$, $\omega \in \Omega$. Indeed, let us prove that the convergence holds in all senses.

Almost Sure Convergence: This is trivial, Calc 1: for a fixed $\omega \in \Omega$,

$$\lim_{n \rightarrow \infty} X_n(\omega) = \lim_{n \rightarrow \infty} \left(\omega + \frac{1}{n} \right)^2 = \omega^2 = X(\omega).$$

So even we have *sure* convergence: the convergence holds for *any* ω , $\{X_n \rightarrow X\} = \Omega$, so $\mathbb{P}(X_n \rightarrow X) = \mathbb{P}(\Omega) = 1$.

Quadratic Mean Convergence: We calculate

$$\mathbb{E}((X_n - X)^2) = \mathbb{E}\left(\left[\left(\omega + \frac{1}{n}\right)^2 - \omega^2\right]^2\right) = \mathbb{E}\left(\left[\frac{2\omega}{n} + \frac{1}{n^2}\right]^2\right)$$

Here we can calculate the expectation in several ways, see later calculations in other examples, but I want to do another trick: since $\omega \in [0, 1]$, then

$$\left[\frac{2\omega}{n} + \frac{1}{n^2}\right]^2 \leq \left[\frac{2}{n} + \frac{1}{n^2}\right]^2 \leq \left[\frac{2}{n} + \frac{1}{n}\right]^2 = \frac{9}{n^2}.$$

Then, by the monotonicity of expectation (that is, of $X \leq Y$ a.s., then $\mathbb{E}(X) \leq \mathbb{E}(Y)$),

$$0 \leq \mathbb{E}((X_n - X)^2) = \mathbb{E}\left(\left[\frac{2\omega}{n} + \frac{1}{n^2}\right]^2\right) \leq \frac{9}{n^2} \rightarrow 0.$$

Then, using the squeezing theorem, we will obtain $\mathbb{E}((X_n - X)^2) \rightarrow 0$.

Remark: In this case, we could calculate the expectation using the formula:

$$\mathbb{E}((X_n - X)^2) = \mathbb{E}\left(\left[\frac{2\omega}{n} + \frac{1}{n^2}\right]^2\right) = \int_0^1 \left[\frac{2\omega}{n} + \frac{1}{n^2}\right]^2 d\mathbb{P} = \int_0^1 \left[\frac{2\omega}{n} + \frac{1}{n^2}\right]^2 d\omega,$$

which is easy to calculate. Please find this formula in advanced Probability textbooks! See also another example below.

Convergence in Probability: We calculate, fixing $\varepsilon > 0$,

$$\mathbb{P}(|X_n - X| < \varepsilon) = \mathbb{P}\left(\left|\frac{2\omega}{n} + \frac{1}{n^2}\right| < \varepsilon\right) = \mathbb{P}\left(\omega \in [0, 1] : \frac{2\omega}{n} < \varepsilon - \frac{1}{n^2}\right) = \mathbb{P}\left(\omega \in [0, 1] : \omega < \frac{n\varepsilon}{2} - \frac{1}{2n}\right).$$

If n is large, then $\frac{n\varepsilon}{2} - \frac{1}{2n} > 1$, so for that large n -s we will have

$$\mathbb{P}(|X_n - X| < \varepsilon) = \mathbb{P}\left(\omega \in [0, 1] : \omega < \frac{n\varepsilon}{2} - \frac{1}{2n}\right) = \mathbb{P}([0, 1]) = 1 \rightarrow 1,$$

as $n \rightarrow \infty$.

Convergence in Distributions: Let us calculate the CDF of X_n first:

$$\begin{aligned} F_{X_n}(x) &= \mathbb{P}(X_n \leq x) = \mathbb{P}(\{\omega \in [0, 1] : X_n(\omega) \leq x\}) = \mathbb{P}(\{\omega \in [0, 1] : \left(\omega + \frac{1}{n}\right)^2 \leq x\}) = \\ &= \begin{cases} 0, & x < 0 \\ \mathbb{P}(\{\omega \in [0, 1] : \omega + \frac{1}{n} \leq \sqrt{x}\}), & x \geq 0 \end{cases} = \begin{cases} 0, & x < 0 \\ \mathbb{P}(\{\omega \in [0, 1] : \omega \leq \sqrt{x} - \frac{1}{n}\}), & x \geq 0 \end{cases} \\ &= \begin{cases} 0, & x < 0 \\ 0, & 0 \leq \sqrt{x} < \frac{1}{n} \\ \mathbb{P}([0, \sqrt{x} - \frac{1}{n}]), & 0 \leq \sqrt{x} - \frac{1}{n} \leq 1 \\ 1, & \sqrt{x} - \frac{1}{n} > 1 \end{cases} = \begin{cases} 0, & x < \frac{1}{n^2} \\ \sqrt{x} - \frac{1}{n}, & \frac{1}{n^2} \leq x \leq \left(1 + \frac{1}{n}\right)^2 \\ 1, & x > \left(1 + \frac{1}{n}\right)^2 \end{cases} \end{aligned}$$

In a similar fashion, $F_X(x)$ is equal to

$$F_X(x) = \begin{cases} 0, & x < 0 \\ \sqrt{x}, & 0 \leq x \leq 1 \\ 1, & x > 1 \end{cases}$$

This F_X is continuous everywhere, so we need to check that $F_{X_n}(x) \rightarrow F_X(x)$ for any $x \in \mathbb{R}$. Clearly, this holds for $x \leq 0$, since $F_{X_n}(x) = F_X(x) = 0$ in that case. If, say $0 < x < 1$, then starting from some number, we will have $\frac{1}{n^2} \leq x \leq \left(1 + \frac{1}{n}\right)^2$, so we will have for large n -s

$$F_{X_n}(x) = \sqrt{x} - \frac{1}{n} \rightarrow \sqrt{x}.$$

In a similar way, one can consider the case $x \geq 1$.

EXAMPLE, CONVERGENCE IN ALL SENSES: Consider again explicitly given r.v. sequence example: the Sample Space is $\Omega = [0, 1]$, with usual length of the interval as a probability, i.e., $\mathbb{P}([a, b]) = b - a$, and consider the following sequence of r.v.:

$$X_n(\omega) = (1 - \omega)^n, \quad \omega \in \Omega, \quad n \in \mathbb{N}.$$

Almost Sure Convergence: Everybody who knows Calc 1 will readily state that

$$\lim_{n \rightarrow \infty} X_n(\omega) = \lim_{n \rightarrow \infty} (1 - \omega)^n = \begin{cases} 0, & \omega \in (0, 1] \\ 1, & \omega = 0. \end{cases}$$

At this point, we could take the right-hand side as the limit of $X_n(\omega)$ and denote by $X(\omega)$, and the proceed to prove that the convergence hold also in all other senses, besides this a.s. (in fact, sure) convergence. But probabilists and real analysts do the following: take $X(\omega) \equiv 0$. This function and the limit of X_n differ only at one point, at $\omega = 0$, so in the probability of one point is zero (recall that the probability of a set is the length of that set). This means that $X_n \rightarrow X$ almost surely. So everybody is happy with this new X - it is much simple than the actual limit, and $X_n \rightarrow X$ a.s. .

Quadratic Mean Convergence: Let us prove also that $X_n \xrightarrow{L^2} X$. To that end, we need to prove that

$$\mathbb{E}\left((X_n - X)^2\right) \rightarrow 0, \quad n \rightarrow \infty,$$

so we need to calculate the expectation $\mathbb{E}\left((X_n - X)^2\right) = \mathbb{E}\left((X_n)^2\right)$. Our tool to calculate this is to find the PDF of X_n and then use it for calculations. So let us find the PDF of X_n .

First, we calculate the CDF of X_n (we need this also for the Convergence in Distributions study):

$$\begin{aligned} F_{X_n}(x) &= \mathbb{P}(X_n \leq x) = \mathbb{P}(\omega \in [0, 1] : (1 - \omega)^n \leq x) = \begin{cases} 0, & x < 0 \\ \mathbb{P}(\omega \in [0, 1] : 1 - \omega \leq \sqrt[n]{x}), & x \geq 0 \end{cases} = \\ &= \begin{cases} 0, & x < 0 \\ \mathbb{P}(\omega \in [0, 1] : \omega \geq 1 - \sqrt[n]{x}), & x \geq 0 \end{cases} = \begin{cases} 0, & x < 0 \\ \mathbb{P}([1 - \sqrt[n]{x}, 1]), & 0 \leq x \leq 1 \\ \mathbb{P}([0, 1]), & x > 1 \end{cases} = \begin{cases} 0, & x < 0 \\ \sqrt[n]{x}, & 0 \leq x \leq 1 \\ 1, & x > 1 \end{cases} \quad (8.1) \end{aligned}$$

Now, the PDF of X_n will be

$$f_{X_n}(x) = F'_{X_n}(x) = \begin{cases} \frac{1}{n} x^{\frac{1}{n}-1}, & 0 < x < 1 \\ 0, & \text{otherwise.} \end{cases}$$

Then,

$$\mathbb{E}\left((X_n - X)^2\right) = \mathbb{E}\left((X_n)^2\right) = \int_{-\infty}^{+\infty} x^2 f_{X_n}(x) dx = \frac{1}{n} \cdot \int_0^1 x^2 \cdot x^{\frac{1}{n}-1} dx = \frac{1}{2n+1} \rightarrow 0.$$

Ura, dami i gospoda!

Remark, somewhat important: We haven't used the following formula to calculate the Expectation, but this is, actually, the definition of the Expectation:

$$\mathbb{E}\left((X_n)^2\right) = \mathbb{E}\left((1 - \omega)^{2n}\right) = \int_0^1 (1 - \omega)^{2n} d\mathbb{P} = \int_0^1 (1 - \omega)^{2n} d\omega = - \int_0^1 (1 - \omega)^{2n} d(1 - \omega) = \frac{1}{2n+1}.$$

See some advanced Probability textbooks!

Convergence in Probability: You are correctly guessing that we are going to torture ourselves by proving that $X_n \xrightarrow{\mathbb{P}} X$, without using the fact that the convergence a.s. implies this.

Assume $\varepsilon > 0$ is fixed. We want to calculate

$$\begin{aligned} \mathbb{P}\left(|X_n - X| < \varepsilon\right) &= \mathbb{P}\left(|X_n| < \varepsilon\right) = \mathbb{P}\left(\omega \in [0, 1] : (1 - \omega)^n < \varepsilon\right) = \mathbb{P}\left(\omega \in [0, 1] : (1 - \omega) < \sqrt[n]{\varepsilon}\right) = \\ &= \mathbb{P}\left(\omega \in [0, 1] : \omega > 1 - \sqrt[n]{\varepsilon}\right) = \begin{cases} \mathbb{P}([0, 1]), & \varepsilon > 1 \\ \mathbb{P}((1 - \sqrt[n]{\varepsilon}, 1]), & 0 < \varepsilon \leq 1 \end{cases} = \begin{cases} 1, & \varepsilon > 1 \\ \sqrt[n]{\varepsilon}, & 0 < \varepsilon \leq 1 \end{cases} \end{aligned}$$

which tends to 1 as $n \rightarrow \infty$, for any fixed $\varepsilon > 0$. This means that $\mathbb{P}\left(|X_n - X| \geq \varepsilon\right) \rightarrow 0$, so X_n tends to X in Probability.

Convergence in Distribution: We have calculated above the CDF F_{X_n} , and from our previous examples we know the value of the CDF F_X . We need to prove that $F_{X_n}(x) \rightarrow F_X(x)$ as $n \rightarrow \infty$, for any $x \neq 0$, for any continuity point of F_X .

From (8.1), this is clear for any $x \notin [0, 1]$. Also, if $x \in (0, 1]$, then

$$F_{X_n}(x) = \sqrt[n]{x} \rightarrow 1 = F_X(x), \quad n \rightarrow \infty.$$

See the Fig. 8.5-8.8, and compare with the graph of the limit 8.1.




Fig. 8.5: Graph of the CDF of X_1 , $F_{X_1}(x)$

EXAMPLE, CONVERGENCE IN DIFFERENT SENSES: Consider again our old friends $\Omega = [0, 1]$, $\mathbb{P}([a, b]) = b - a$ for $[a, b] \subset [0, 1]$, and define

$$X_n(\omega) = \begin{cases} 0, & \omega \in [0, 1] \setminus \left\{ \frac{1}{n}, \frac{2}{n}, \dots, \frac{n-1}{n}, 1 \right\} \\ 1, & \omega \in \left\{ \frac{1}{n}, \frac{2}{n}, \dots, \frac{n-1}{n}, 1 \right\}. \end{cases}$$

Say, X_3 is equal to 1 at $\frac{1}{3}, \frac{2}{3}$ and $\frac{3}{3} = 1$, and is equal to 0 at any other point of $[0, 1]$. X_4 is equal to 1 at $\frac{1}{4}, \frac{2}{4}, \frac{3}{4}$ and $\frac{4}{4} = 1$, and is equal to 0 at any other point and so on.

Convergence in Almost Sure sense: We can prove that $X_n \rightarrow X$ a.s., where $X(\omega) \equiv 0$. Indeed, if $\omega \in [0, 1]$ is irrational, then $X_n(\omega) = 0 \rightarrow 0 = X(\omega)$. Now, $X_n \not\rightarrow X$ can happen only at rational points, but the length (probability) of the set of all rational points from $[0, 1]$ is 0. So

$$\mathbb{P}(X_n \rightarrow X) = 1.$$




Fig. 8.6: Graph of the CDF of X_2 , $F_{X_2}(x)$ 

Fig. 8.7: Graph of the CDF of X_5 , $F_{X_5}(x)$. In fact, here, and in the next figure, you see a jump at $x = 0$, but, in fact, the function is continuous. This is a drawback of computer software - it cannot give very high precision in this example to draw close to 0. Moral: do not trust computers much! Do not trust man either. Trust only good old friends. Say, Math 😊

Convergence in Probability: We fix $\varepsilon > 0$, and calculate

$$\mathbb{P}(|X_n - X| \geq \varepsilon) = \mathbb{P}(X_n \geq \varepsilon) = \begin{cases} 0, & \varepsilon > 1 \\ \mathbb{P}(X_n = 1), & \varepsilon \leq 1 \end{cases} = 0$$

Fig. 8.8: Graph of the CDF of X_{15} , $F_{X_{15}}(x)$

Convergence in Quadratic Mean: We calculate $\mathbb{E}((X_n - X)^2) = \mathbb{E}(X_n^2)$ by considering that X_n is a discrete r.v. with value 0 with probability 1: $\mathbb{P}(X_n = 0) = 1$. So by the discrete r.v. expectation formula, $\mathbb{E}(X_n^2) = 0^2 \cdot \mathbb{P}(X_n = 0) = 0$.

We could do this by calculating first the CDF of X_n , F_{X_n} :

$$F_{X_n}(x) = \mathbb{P}(X_n \leq x) = \begin{cases} 0, & x < 0 \\ \mathbb{P}(X_n = 0), & 0 \leq x < 1 \\ 1, & x \geq 1 \end{cases} = \begin{cases} 0, & x < 0 \\ 1, & x \geq 0 \end{cases}$$

This means that X_n is a discrete r.v. with the value 0 with probability 1: $\mathbb{P}(X_n = 0) = 1$. The rest is as above.

Convergence in Distributions: This is simple, since, as we have calculated above,

$$F_{X_n}(x) = \begin{cases} 0, & x < 0 \\ 1, & x \geq 0 \end{cases} = F_X(x) \quad x \in \mathbb{R}.$$

REMARK, TACTICS TO PROVE THE CONVERGENCE OF A R.V. SEQUENCE, EXPLICIT CASE: In the case when we have our sequence $X_n(\omega)$ given explicitly, as a function of ω , then we can try to prove first the a.s. convergence. To that end, we consider $X_n(\omega)$ for a fixed ω and try to calculate its limit. This is the same as in the Real Analysis course we were calculating functional sequence pointwise limit $\lim f_n(x)$ or in the Calculus class we were calculating limits with parameters, say, the limit of x^n , when $n \rightarrow \infty$. If you obtain that $X_n \rightarrow X$ a.s. for some X , then the convergence hold also in the Probability and Distribution senses.

8.2.2 R.V. described through their Distribution

EXAMPLE, CONVERGENCE OF A SEQUENCE GIVEN BY THEIR DISTRIBUTIONS: First, very simple (and almost not random) example: assume $X_n = \frac{1}{n}$ with probability 1. Then $X_n \rightarrow 0$ in Distribution. If all X_n -s are defined on the same Probability Space, then also $X_n \rightarrow 0$ in Quadratic Mean and in Probability.

We can prove also the a.s. convergence. If we will denote by A_n the set, where $X_n \neq \frac{1}{n}$. By the condition above, $\mathbb{P}(A_n) = 0$. Denote now $A = \bigcup_{n=1}^{\infty} A_n$. Then $\mathbb{P}(A) = 0$ too. And in the complement of A_n , we will have that $X_n(\omega) = \frac{1}{n}$. Now, for any x of this type, we will have that $X_n(\omega) \rightarrow X(\omega)$ a.s.

EXAMPLE, CONVERGENCE OF A SEQUENCE GIVEN BY THEIR DISTRIBUTIONS: We consider the following experiment: we are tossing a coin infinitely many times (countably infinite, of course). Let X_n be 1, if n -th toss resulted in heads, and let X_n be 0 otherwise⁸.

Now, is X_n a.s. convergent? Let us find the set of all points where X_n converges. If for some scenario X_n converges, then in that scenario either starting from some point on we will have heads shown, or we will have only tails shown⁹. And it can be shown that the set of all scenarios when X_n converges, is a null set (set of probability 0), so our X_n will not converge a.s..

Also, it will not converge in Probability (prove this!). But since all $X_n \sim \text{Bernoulli}(0.5)$, then X_n will tend in Distribution to a r.v. X with $X \sim \text{Bernoulli}(0.5)$. ■

EXAMPLE, CONVERGENCE IN DISTRIBUTION: Let $n \in \mathbb{N}$. Assume X_n is a discrete r.v. with

$$X_n \sim \text{DiscreteUnif}\left(n; \frac{1}{n}, \frac{2}{n}, \frac{3}{n}, \dots, 1\right),$$

i.e.,

Values of X_n	$\frac{1}{n}$	$\frac{2}{n}$	$\frac{3}{n}$...	1
PMF of X_n	$\frac{1}{n}$	$\frac{1}{n}$	$\frac{1}{n}$...	$\frac{1}{n}$.

In fact, here we are not specifying the underlying probability space, so the r.v.'s X_n can be defined on different probability spaces.

Intuitively, the sequence X_n will tend in Distributions to the Standard Uniform Distributed r.v.. if $X \sim \text{Unif}([0, 1])$, then it is easy to prove that $X_n \xrightarrow{D} X$. Say, you can run the R code given below to see the picture.

In this example, we cannot talk about the a.s. or in Probability convergence of X_n , since they can be defined on different probability spaces.

R CODE, PREVIOUS EXAMPLE CODE: Here we want to draw the above distribution CDFs.

```
par(mfcol = c(1,2))
```

```
n <- 20
```

```
x <- (1:n)/n
```

```

f <- ecdf(x)
plot(f, xlim = c(-0.5,1.5), lwd = 2, main = "n = 20", ylab = "CDF")
#Now the Uniform
par(new = T)
plot(punif, xlim = c(-0.5,1.5), col = "red", lwd = 3, main = "n = 20", ylab = "CDF")
par(new = F)

n <- 50
x <- (1:n)/n
f <- ecdf(x)
plot(f, xlim = c(-0.5,1.5), lwd = 2, main = "n = 50", ylab = "CDF")
#Now the Uniform
par(new = T)
plot(punif, xlim = c(-0.5,1.5), col = "red", lwd = 3, main = "n = 50", ylab = "CDF")
par(new = F)

```

EXAMPLE, CONVERGENCE IN PROBABILITY AND DISTRIBUTIONS: Assume X_n is a Discrete r.v. with the following PMF:

$$\begin{array}{c|c|c} X_n & 5 - \frac{1}{n^2} & n \\ \hline \mathbb{P}(X_n = x) & 1 - \frac{1}{n} & \frac{1}{n}. \end{array}$$

Then $X_n \xrightarrow{\mathbb{P}} X$ and $X_n \xrightarrow{D} X$, where $X \equiv 5$ (of course, the last statement follows from the former one, but let us prove independently).

Remark: When dealing with the Convergence in Probability, we need to assume that all X_n -s are defined on the same Probability Space. In fact, since X , the limiting r.v., is constant, then we can calculate $\mathbb{P}(|X_n - X| \geq \varepsilon)$ even if X_n -s are defined on different Probability Spaces.

Convergence in Probability: We fix $\varepsilon > 0$ and calculate the probability $\mathbb{P}(|X_n - X| \geq \varepsilon)$. To that end, we can see that the r.v. $|X_n - X|$ will take the values $\frac{1}{n^2}$ and $|n - 5|$ with probabilities, $1 - \frac{1}{n}$ and $\frac{1}{n}$, respectively, i.e. $|X_n - X|$ has the distribution

$$\begin{array}{c|c|c} |X_n - X| & \frac{1}{n^2} & |n - 5| \\ \hline \mathbb{P}(|X_n - X| = x) & 1 - \frac{1}{n} & \frac{1}{n}. \end{array}$$

We want to see when we can have $|X_n - X| \geq \varepsilon$. If $n > \frac{1}{\sqrt{\varepsilon}}$, then we will have $\frac{1}{n^2} < \varepsilon$, so $|X_n - X| \geq \varepsilon$ can happen only if $|X_n - X| = |n - 5|$, and $|n - 5| \geq \varepsilon$, and the probability of that is $\frac{1}{n}$. This means that

$$\mathbb{P}(|X_n - X| \geq \varepsilon) \stackrel{n > \frac{1}{\sqrt{\varepsilon}}, |n-5| \geq \varepsilon}{=} \mathbb{P}(|X_n - X| = |n - 5|) = \frac{1}{n} \rightarrow 0.$$

Convergence in Distribution: The CDF of X is

$$F_X(x) = \begin{cases} 0, & x < 5 \\ 1, & x \geq 5 \end{cases}$$

The only discontinuity point of F_X is $x = 5$. The CDF of X_n is, for $n \geq 5$,

$$F_{X_n}(x) = \begin{cases} 0, & x < 5 - \frac{1}{n^2} \\ 1 - \frac{1}{n}, & 5 - \frac{1}{n^2} \leq x < n \\ 1, & x \geq n. \end{cases}$$

We need to prove that $F_{X_n}(x) \rightarrow F_X(x)$ for any $x \neq 5$. Say, $x > 5$. Then, starting from some number on, we will have $5 - \frac{1}{n^2} \leq x < n$, i.e.

$$F_{X_n}(x) = 1 - \frac{1}{n} \rightarrow 1 = F_X(x).$$

If $x < 5$, then starting from some number on, we will have $x < 5 - \frac{1}{n^2}$, and then

$$F_{X_n}(x) = 0 \rightarrow 0 = F_X(x).$$

Try to draw the graphs to see the convergence visually!

Remark: Note that in this case, we will not have a convergence in Quadratic Mean. This is because the distribution of $(X_n - X)^2$ will be

$$\frac{(X_n - X)^2}{\mathbb{P}((X_n - X)^2 = x)} \parallel \begin{array}{|c|} \hline \frac{1}{n^4} \\ \hline \end{array} \begin{array}{|c|} \hline (n-5)^2 \\ \hline \end{array}$$

$$\parallel \begin{array}{|c|} \hline 1 - \frac{1}{n} \\ \hline \end{array} \begin{array}{|c|} \hline \frac{1}{n} \\ \hline \end{array}$$

and

$$\mathbb{E}((X_n - X)^2) = \left(1 - \frac{1}{n}\right) \cdot \frac{1}{n^4} + \frac{1}{n} \cdot (n-5)^2 \rightarrow +\infty.$$

EXAMPLE, CONVERGENCE IN PROBABILITY AND DISTRIBUTIONS: Assume $X_n \sim \mathcal{N}\left(0, \frac{1}{n}\right)$. Then $X_n \xrightarrow{D} 0$. ' In general, we cannot talk about the convergence in Probability or a.s. or L^2 means in this case, since X_n -s can be r.v.s defined in very different sample spaces. If we will assume that they are defined on the same Probability Space, then one can prove that $X_n \xrightarrow{P} 0$, and also $X_n \xrightarrow{q.m.} 0$.

EXAMPLE, CONVERGENCE IN PROBABILITY AND DISTRIBUTIONS: Assume $X_n \sim \text{Unif}\left[0, \frac{1}{n}\right]$. Then $X_n \xrightarrow{D} 0$. If all X_n are defined on the same Ω , say, $\Omega = [0, 1]$, then also $X_n \xrightarrow{P} 0$ and $X_n \xrightarrow{q.m.} 0$.

EXAMPLE, POISSON AS A LIMIT OF BINOMIAL: Assume $X_n \sim \text{Binomial}(n, \frac{\lambda}{n})$ and $X \sim \text{Pois}(\lambda)$. Then $X_n \xrightarrow{D} X$.

REMARK, CONVERGENCE IN PROBABILITY: As we have noted above, except the case of the convergence in Distributions, for all other three types of convergence we need to have that all X_k -s and X are defined on the same Probability Space. This is to ensure that we can calculate, say $\mathbb{E}((X_n - X)^2)$ or $\mathbb{P}(|X_n - X| \geq \varepsilon)$. But, in fact, there is one particular case, when we can talk about the convergence of a sequence of r.v. X_k to X in the L^2 or in the Probability sense even when X_k and X are not defined on the same probability space. This case is when $X \equiv \text{const}$. In that case, for any $n \in \mathbb{N}$, we can calculate

$$\mathbb{E}((X_n - X)^2) \quad \text{and} \quad \mathbb{P}(|X_n - X| \geq \varepsilon)$$

and see if these numerical sequences tend to 0 or not.

8.3 The LLN and CLT

Assume now X_n is a sequence of IID r.v. . This means that all X_k -s have the same distribution. In particular, this means that if one of X_k -s has a finite expectation and variance, then all X_k -s have finite expectations and variances and

$$\mathbb{E}(X_i) = \mathbb{E}(X_j), \quad \text{and} \quad \text{Var}(X_i) = \text{Var}(X_j), \quad \forall i, j \in \mathbb{N}.$$

The Idea of LLN and CLT In Probability Theory and Statistics, one of the important questions is to study the distribution and/or the behavior of either the sum

$$S_n = X_1 + X_2 + \dots + X_n$$

or the average

$$\bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n}.$$

In fact, these problems are not easy ones. We know from the Probability Theory that the PDF of $X + Y$ in case when X and Y are independent, is the convolution of the PDFs of X and Y . And calculation of the convolution is not an easy task. And, for S_n or X_n , we need to calculate $n - 1$ convolutions!

Let us see what we can say about S_n and \bar{X}_n , in general. Since X_k -s are IID, they have the same mean and variance, and let

$$\mathbb{E}(X_k) = \mu \quad \text{and} \quad \text{Var}(X_k) = \sigma^2.$$

Proposition 8.5. *If X_k are IID with the above expectation and variance, then¹⁰*

$$\begin{aligned} \mathbb{E}(S_n) &= n \cdot \mu, & \mathbb{E}(\bar{X}_n) &= \mu; \\ \text{Var}(S_n) &= n \cdot \sigma^2, & \text{Var}(\bar{X}_n) &= \frac{\sigma^2}{n}. \end{aligned}$$

Proof. Obvious, we will meet this many times! □

Say, what information is giving this proposition about \bar{X}_n ? If n is large, then the possible values of \bar{X}_n are around μ , since $\mathbb{E}(\bar{X}_n) = \mu$, and are very concentrated around μ , since $\text{Var}(\bar{X}_n) = \frac{\sigma^2}{n}$ is small (n is large).

But, of course, having only the expectation and variance of S_n and X_n , is giving us just a very partial information about their distributions. Sometimes, for some particular cases, we can exactly find the distributions of S_n and/or \bar{X}_n .

¹⁰We can read the assertion $\mathbb{E}(X_n) = \mu$ as *the mean of the means is the mean* ☺

Proposition 8.6. a. If $X_k \sim \mathcal{N}(\mu, \sigma^2)$, $k = 1, \dots, n$, are independent, then

$$S_n = X_1 + \dots + X_n \sim \mathcal{N}(n \cdot \mu, n \cdot \sigma^2) \quad \text{and} \quad \bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n} \sim \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right).$$

b. If $X_k \sim \text{Bernoulli}(p)$, $k = 1, \dots, n$ are independent, then

$$S_n = X_1 + \dots + X_n \sim \text{Binom}(n, p);$$

c. If $X_k \sim \text{Binom}(m, p)$, $k = 1, \dots, n$, are independent, then¹¹

$$S_n = X_1 + \dots + X_n \sim \text{Binom}(n \cdot m, p).$$

d. If $X_k \sim \text{Geom}(p)$, $k = 1, \dots, n$, are independent, then¹²

$$S_n = X_1 + \dots + X_n \sim \text{NegBin}(n, p)$$

e. If $X_k \sim \text{Pois}(\lambda)$, $k = 1, \dots, n$, are independent, then

$$S_n = X_1 + \dots + X_n \sim \text{Pois}(n \cdot \lambda).$$

f. If $X_k \sim \text{Gamma}(\alpha, \beta)$, $k = 1, \dots, n$, are independent, then

$$S_n = X_1 + \dots + X_n \sim \text{Gamma}(n \cdot \alpha, \beta).$$

In particular, if $X_k \sim \text{Exp}(\lambda)$, $k = 1, \dots, n$, are independent, then

$$S_n = X_1 + \dots + X_n \sim \text{Gamma}(n, \lambda).$$

But, unfortunately, the distribution of S_n or \bar{X}_n cannot be calculated in simple terms for many cases. And here we can use the LLN and/or CLT. The Law of Large Numbers (LLN) is giving the asymptotic behavior of \bar{X}_n , and the Central Limit Theorem describes the asymptotic distribution of S_n and \bar{X}_n .

Theorem 8.3 (The Strong Law of Large Numbers). If X_1, \dots, X_n, \dots is a sequence of IID r.v. with $\mathbb{E}(|X_1|) < +\infty$ and if $\mathbb{E}(X_i) = \mu$, then

$$\bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n} \rightarrow \mu \quad \text{a.s.}$$

Theorem 8.4 (The Weak Law of Large Numbers). If X_1, \dots, X_n, \dots is a sequence of IID r.v. with finite expectation $\mathbb{E}(X_i) = \mu$ and variance $\text{Var}(X_i) = \sigma^2$, then

$$\bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n} \xrightarrow{P} \mu.$$

¹¹The distribution of \bar{X}_n can be described, but is not of standard ones. Just divide the values of S_n by n .

¹²The Negative Binomial Distribution is the number of failures before the n -th success when doing Bernoulli(p) trials, see https://en.wikipedia.org/wiki/Negative_binomial_distribution.

R CODE, LLN:

```
#LLN
n <- 1000 #number of r.v.'s
expect <- 0.6 #this will be the expectation of each random variable
X <- rbinom(n, 1, expect) #generating n samples from the same distribution
S <- cumsum(X) #calculating the cumulative sum: S = (X_1, X_1+X_2, X_1+X_2+X_3,...)
p <- S/(1:n) #This will produce p=(X_1/1, (X_1+X_2)/2, (X_1+X_2+X_3)/3,...)
plot(p, type = "l")
abline(expect,0, col = "red", lwd = 2) #giving in red the limit
```

The result is given in Fig. 8.9.




Fig. 8.9: The Law of Large Numbers

REMARK, USING LLN TO CALCULATE PROBABILITIES BY SIMULATIONS: Great and interesting thing is that we can use the LLN to calculate probabilities, areas, volumes, one and multi-dimensional integrals by Simulations. The method is called the Monte-Carlo Method, and is widely used in many applications.

To describe the ideas, assume we want to calculate the probability $\mathbb{P}(a \leq X \leq b)$ for some continuous r.v. X . Assume we can generate random numbers from the distribution of X . How to use this to calculate the mentioned probability?

We consider the indicator function $\mathbb{1}_{[a,b]}(x)$. If we will consider the r.v. $\mathbb{1}_{[a,b]}(X)$, and calculate the expected value $\mathbb{E}(\mathbb{1}_{[a,b]}(X))$, then we will get

$$\mathbb{E}(\mathbb{1}_{[a,b]}(X)) = \int_{-\infty}^{+\infty} \mathbb{1}_{[a,b]}(x) f_X(x) dx = \int_a^b f(x) dx = \mathbb{P}(a \leq X \leq b).$$

So we need to calculate the expected value $\mathbb{E}(\mathbb{1}_{[a,b]}(X))$. And we can approximate this by the LLN! Ura! Here is the method: by the LLN, if we will take r.v.s X_1, X_2, \dots, X_n that are IID with the same

distribution as X , then

$$\frac{\mathbb{1}_{[a,b]}(X_1) + \mathbb{1}_{[a,b]}(X_2) + \dots + \mathbb{1}_{[a,b]}(X_n)}{n} \rightarrow \mathbb{E}(\mathbb{1}_{[a,b]}(X)).$$

, so

$$\mathbb{E}(\mathbb{1}_{[a,b]}(X)) \approx \frac{\mathbb{1}_{[a,b]}(X_1) + \mathbb{1}_{[a,b]}(X_2) + \dots + \mathbb{1}_{[a,b]}(X_n)}{n}.$$

The numerator in the fraction on the RHS is just the number of X_k -s in $[a, b]$. So basically,

$$\mathbb{P}(a \leq X \leq b) = \mathbb{E}(\mathbb{1}_{[a,b]}(X)) \approx \frac{\# \text{ of } X_k \text{ in } [a, b]}{n}.$$

And the code will go as follows: we generate x_1, \dots, x_n from the distribution of X , calculate the number of x_k -s in $[a, b]$, and divide this number to the total number of elements generated, n . This ratio will give an approximation for the probability $\mathbb{P}(a \leq X \leq b)$. Here is the code:

R CODE, USING LLN TO CALCULATE PROBABILITIES BY SIMULATIONS: Assume $X \sim \mathcal{N}(0, 1)$, and we want to calculate (approximately) the probability

$$\mathbb{P}(0.3 \leq X \leq 1.4).$$

The actual probability is equal to

$$\mathbb{P}(0.3 \leq X \leq 1.4) = \frac{1}{\sqrt{2\pi}} \cdot \int_{0.3}^{1.4} e^{-x^2/2} dx,$$

and calculation of the last integral cannot be done easily - we need to use some numerical technique. Here we will use the idea above:

```
#Probabilities by Simulations
n <- 5000
a <- 0.3
b <- 1.4
x <- rnorm(n)
inside <- sum((x<=b) & (x>=a))
prob <- inside/n
real.prob <- pnorm(b)-pnorm(a)
```

I love Math! Well, of course, there is some hidden thing here - we need to be able to generate normal random variables. But since **R** is doing this for us, we can happily use this opportunity to calculate some complicated integrals 😊

REMARK, USING LLN TO CALCULATE INTEGRALS: As we have mentioned above, we can use the LLN and Monte Carlo Simulations to calculate integrals. Let me explain ideas for a 1D integral calculation. Assume we have a positive continuous function $g : [a, b] \rightarrow \mathbb{R}$, and we want to calculate approximately

$$\int_a^b g(x) dx.$$

Nothing probabilistic yet. But let us make this probabilistic, let us write this integral as an expected value for some r.v.. To that end, consider a r.v. $X \sim \text{Unif}[a, b]$. Then the PDF of X will be $f_X(x) = \frac{1}{b-a}$ for $x \in [a, b]$ and $f_X(x) = 0$, for $x \notin [a, b]$. Now, consider the r.v. $g(X)$. We know that

$$\mathbb{E}(g(X)) = \int_{-\infty}^{+\infty} g(x)f_X(x)dx = \frac{1}{b-a} \cdot \int_a^b g(x)dx,$$

so

$$\int_a^b g(x)dx = (b-a) \cdot \mathbb{E}(g(X)).$$

Hopplya!! Magic! Non-probabilistic, our old and good friend Calculus integral is an expected value of some random variable! Now, to calculate the expectation, we use the LLN: we take IID r.v. $X_1, X_2, \dots, X_n \sim \text{Unif}[a, b]$, and then

$$\frac{g(X_1) + g(X_2) + \dots + g(X_n)}{n} \rightarrow \mathbb{E}(g(X)),$$

so

$$\int_a^b g(x)dx = (b-a) \cdot \mathbb{E}(g(X)) \approx (b-a) \cdot \frac{g(X_1) + g(X_2) + \dots + g(X_n)}{n}.$$

Well, maybe one is not using this method for 1D or 2D integrals, because in this cases we have well-developed theory for numerical calculation of integrals, and many nice methods, with proven error estimation and analysis. But for multidimensional integrals, this simulations method will give some good results compared to other numerical methods.

The above method is implemented below.

R CODE, LLN FOR INTEGRAL CALCULATION:

```
#Integral by Simulations
#Say, we want to integrate the function x^2*sin(x^4) over [0,3]
g <- function(x){
  return(x^2*sin(x^4))
}
a <- 0
b <- 3
n <- 10000
x <- runif(n, min = a, max = b)
approx_int <- (b-a)*mean(g(x))
int <- integrate(g,a,b) #R-s native integral calculator
abs(int$value - approx_int) #absolute difference between R-s native integral calculator's value
```

REMARK, LLN FOR INTEGRAL CALCULATION, 2D VERSION: Explain here how to calculate

$$\int_a^b g(x)dx$$

by using 2D ideas - calculation of the area under the curve $y = g(x)$ in $x \in [a, b]$.

REMARK, LLN, AGAIN: Assume X_k 's are IID with the mean $\mathbb{E}(X_k) = \mu$ and $\text{Var}(X_k) = \sigma^2$. The LLN states that

$$\frac{X_1 + X_2 + \dots + X_n}{n} \rightarrow \mu$$

(in Probability or in the a.s. sense), and an easy fact from the Proposition 8.5 says that

$$\mathbb{E}\left(\frac{X_1 + X_2 + \dots + X_n}{n}\right) = \mu$$

Do you see what is the difference between the LLN (in some form) and that very easy fact?

The latter one means that the means \bar{X}_n are equal to the mean μ in the mean \smile . So it says that the values of \bar{X}_n are around μ , and if we will generate a lot of values for \bar{X}_n , the average of that values will be close to μ . Here n is fixed, the number of generated values of \bar{X}_n is large.

The former says that **almost all** values of \bar{X}_n are close to μ , if n is large enough.

This theorems state that random variables \bar{X}_n become more and more concentrated around μ , around their mean. To give more accurate information, to give the asymptotic distribution of \bar{X}_n , we use the CLT:

Theorem 8.5 (The Central Limit Theorem). Let X_n be a sequence of iid r.v. with finite expectation $\mu = \mathbb{E}(X_i)$ and variance $\sigma^2 = \text{Var}(X_i)$. Denote

$$\bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n}, \quad n \in \mathbb{N}.$$

Then if we will standardize (normalize)¹³ \bar{X}_n , i.e. if we will denote

$$Z_n = \frac{\bar{X}_n - \mathbb{E}(\bar{X}_n)}{\sqrt{\text{Var}(\bar{X}_n)}} = \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}},$$

then

$$Z_n \xrightarrow{D} Z$$

for some $Z \sim \mathcal{N}(0, 1)$.

REMARK, CLT: The conclusion of the CLT is that

$$\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \xrightarrow{D} Z$$

with $Z \sim \mathcal{N}(0, 1)$, which means that for any $a < b$,

$$\mathbb{P}\left(a \leq \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \leq b\right) \rightarrow \mathbb{P}(a \leq Z \leq b) = \frac{1}{\sqrt{2\pi}} \cdot \int_a^b e^{-x^2/2} dx,$$

¹³If X is a r.v., then by its standardization we mean creating another r.v. by shifting and scaling X , which will have an expectation 0 and variance 1. If we will denote

$$Y = \frac{X - \mathbb{E}(X)}{\text{SD}(X)} = \frac{X - \mathbb{E}(X)}{\sqrt{\text{Var}(X)}},$$

then

$$\mathbb{E}(Y) = 0 \quad \text{and} \quad \text{Var}(Y) = 1.$$

so for large n ,

$$\mathbb{P}(a \leq \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \leq b) \approx \frac{1}{\sqrt{2\pi}} \cdot \int_a^b e^{-x^2/2} dx,$$

In the case when $Z_n \xrightarrow{D} Z$ for some $Z \sim \mathcal{N}(0, 1)$, one uses the notation

$$Z_n \xrightarrow{D} \mathcal{N}(0, 1).$$

So the CLT states that

$$\frac{\sqrt{n} \cdot (\bar{X}_n - \mu)}{\sigma} \xrightarrow{D} \mathcal{N}(0, 1),$$

so

$$\sqrt{n} \cdot (\bar{X}_n - \mu) \xrightarrow{D} \sigma \mathcal{N}(0, 1),$$

or

$$\sqrt{n} \cdot (\bar{X}_n - \mu) \xrightarrow{D} \mathcal{N}(0, \sigma^2).$$

If we will use a little bit not rigorous math terms, then we can write this as

$$\sqrt{n} \cdot (\bar{X}_n - \mu) \approx \mathcal{N}(0, \sigma^2),$$

which further can be written in the form¹⁴

$$\bar{X}_n - \mu \approx \frac{1}{\sqrt{n}} \cdot \mathcal{N}(0, \sigma^2) = \mathcal{N}\left(0, \frac{\sigma^2}{n}\right),$$

and then

$$\bar{X}_n \approx \mu + \mathcal{N}\left(0, \frac{\sigma^2}{n}\right) = \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right).$$

So, CLT states that if $\{X_n\}$ is a sequence of IID r.v., then the asymptotic distribution of \bar{X}_n is

$$\bar{X}_n \approx \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right) \quad \text{for large } n,$$

that is, for large n , the approximate distribution of \bar{X}_n is Normal with parameters μ and $\frac{\sigma^2}{n}$, and this is *independent of the actual distribution of X_k -s!*

REMARK, CLT, IN OTHER FORM: Sometimes people use the CLT not for the sequence \bar{X}_n , but for

$$S_n = X_1 + X_2 + \dots + X_n.$$

The process is just as above: we first standardize S_n : since $\mathbb{E}(S_n) = n \cdot \mathbb{E}(X_1) = n \cdot \mu$, and $\text{Var}(S_n) = n \cdot \text{Var}(X_1) = n\sigma^2$, then the r.v. Z_n defined by

$$Z_n = \frac{S_n - n\mu}{\sqrt{n} \cdot \sigma}$$

will have¹⁵ an expectation 0 and variance 1:

$$\mathbb{E}(Z_n) = 0 \quad \text{and} \quad \text{Var}(Z_n) = 1.$$

¹⁴We write $k \cdot \mathcal{N}(0, \sigma^2) = \mathcal{N}(0, k^2 \sigma^2)$ in the sense that if $X \sim \mathcal{N}(0, \sigma^2)$, then $k \cdot X \sim \mathcal{N}(0, k^2 \cdot \sigma^2)$

Then the CLT will then give:

$$Z_n \xrightarrow{D} \mathcal{N}(0, 1).$$

If we will use some not rigorous notations like above, we can write

$$S_n = n \cdot \bar{X}_n \approx n \cdot \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right) = \mathcal{N}(n \cdot \mu, n \cdot \sigma^2) \quad \text{for large } n,$$

that is,

$$S_n \approx \mathcal{N}(n \cdot \mu, n \cdot \sigma^2) \quad \text{for large } n.$$

This means that the asymptotic distribution of the sum S_n is Normal with the mean $n \cdot \mu$ and variance $n \cdot \sigma^2$.

R CODE, CLT:

```
#The Central Limit Theorem
no_of_simulations <- 5000
sample_size <- 5000
m <- c() # in m we will keep the values of normalized means

#Say, we will generate from Binom(10,0.4)
mu <- 10*0.4 #the expected value for a Binom(10,0.4) r.v
sig <- sqrt(10*0.4*(1-0.4)) #the standard deviation for a Binom(10,0.4) r.v

#Uncomment, if you want to generate from Unif[4,10]
#mu <- (4+10)/2 #The Expected Values for Unif[4,10] r.v
#sig <- sqrt((10-4)^2/12) #the standard deviation for a Unif[4,10] r.v

for(i in 1:no_of_simulations){
  x <- rbinom(sample_size, size = 10, prob = 0.4) # we generate a sample from the Binom(10, 0.4)
  #x <- runif(sample_size, min = 4, max = 10) # we generate a sample from the Unif[4,10]
  m[i] <- (sum(x) - sample_size* mu)/(sqrt(sample_size)*sig)
}
#plotting the results
hist(m, breaks = seq(min(m)-0.2, max(m)+0.2, by = 0.2), col = "lightcyan", freq = F, xlim = c(-3,3))
par(new = T)
curve(dnorm, xlim = c(-3,3), ylim = c(0,0.45), col = "red", lwd = 2)
```

R CODE, CLT, ANOTHER INTERPRETATION: Here we give another method to demonstrate the CLT, and also to calculate probabilities by simulations. The idea is the following: according to CLT, for any a, b we need to have

$$\mathbb{P}\left(a \leq \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \leq b\right) \approx \frac{1}{\sqrt{2\pi}} \cdot \int_a^b e^{-x^2/2} dx,$$

for a large n . We want to check this using **R**.

We will calculate the probability on the LHS using the following idea: we will take an observation x_1, \dots, x_n , calculate the value of \bar{X}_n , and then $\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}}$. And we will repeat this process and obtain

different values for $\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}}$. And to calculate the probability that this quantity is between a and b, we will calculate how many of our generated values are in [a, b], and divide that number to the total number of experiments (simulations). This ratio will be close to the probability above, if the number of simulations is large enough. Here is the code:

```
## Second method to check CLT Result, with probabilities

no_of_simulations <- 1000
sample_size <- 1000
standardized_x <- c()
#we will use gamma-distributed r.v.s
g_shape = 2 #parameters of gamma-distrib
g_scale = 1 #parameters of gamma-distrib
mu <- g_shape*g_scale # the expectation of gamma distrib
sig <- sqrt(g_shape*g_scale^2) # the standard deviation of the gamma distrib

for (i in 1:no_of_simulations){
  x <- rgamma(sample_size, shape = g_shape, scale = g_scale)
  #we will use this time the averages form of the CLT
  standardized_x[i] <- sqrt(sample_size)*(mean(x)-mu)/sig
}

#We want to calculate the probability that the standardized_x is in [a,b]
a <- -0.4
b <- 0.6
N <- sum((standardized_x<=b) & (standardized_x>=a))
prob <- N/no_of_simulations #Probability by Simulations
prob_with_normal <- pnorm(b)-pnorm(a) #Probability calculated by the Standard Normal CDF
prob
prob_with_normal
```

REMARK, PROP, LLN AND CLT: I am in love with the LLN, CLT and beautiful things like that, that's why I want to give the general idea again. This time, for \bar{X}_n only. So what info we have obtained from Proposition 8.5, LLN and CLT about \bar{X}_n ? They are gradually giving us the following info:

Proposition 8.5:

$$\mathbb{E}\left(\frac{X_1 + X_2 + \dots + X_n}{n}\right) = \mu \quad \text{and} \quad \text{Var}\left(\frac{X_1 + X_2 + \dots + X_n}{n}\right) = \frac{\sigma^2}{n};$$

LLN:

$$\frac{X_1 + X_2 + \dots + X_n}{n} \rightarrow \mu$$

in the Probability or in the a.s. sense;

CLT:

$$\frac{X_1 + X_2 + \dots + X_n}{n} \approx \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right) \quad \text{for large } n$$

REMARK, CLT, ONCE AGAIN: Let me give the CLT for both cases simultaneously. Assume X_n be a sequence of IID r.v. with finite expectation $\mu = \mathbb{E}(X_i)$ and variance $\sigma^2 = \text{Var}(X_i)$. The steps of CLT are as follows:

Step	For the sum $S_n = X_1 + X_2 + \dots + X_n$	For the Sample Mean $\bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n} = \frac{S_n}{n}$
1. Standardization (Normalization)	$Z_n = \frac{S_n - \mathbb{E}(S_n)}{\sqrt{\text{Var}(S_n)}}$	$Z_n = \frac{\bar{X}_n - \mathbb{E}(\bar{X}_n)}{\sqrt{\text{Var}(\bar{X}_n)}}$
2. Calculation	$\mathbb{E}(S_n) = n \cdot \mu, \text{Var}(S_n) = n \cdot \sigma^2$	$\mathbb{E}(\bar{X}_n) = \mu, \text{Var}(\bar{X}_n) = \frac{\sigma^2}{n}$
3. The Value of Z_n simplified	$Z_n = \frac{S_n - n \cdot \mu}{\sqrt{n} \cdot \sigma}$	$Z_n = \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}}$
4. Conclusion	$Z_n = \frac{S_n - n \cdot \mu}{\sqrt{n} \cdot \sigma} \xrightarrow{D} \mathcal{N}(0, 1)$	$Z_n = \frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \xrightarrow{D} \mathcal{N}(0, 1)$

I syo!

8.4 Supplement

Let us give examples of Identically Distributed (ID) random variables:

EXAMPLE, ID RANDOM VARIABLES: Let us construct several r.v. with Bernoulli(0.5) distribution.

- The experiment is a one toss of a fair coin. Then $\Omega = \{H, T\}$. Take $X_1(H) = 0$ and $X_1(T) = 1$. Then $X_1 \sim \text{Bernoulli}(0.5)$;
- The experiment is again a one toss of a fair coin. Then $\Omega = \{H, T\}$. Take $X_2(H) = 1$ and $X_2(T) = 0$. Then $X_2 \sim \text{Bernoulli}(0.5)$;
- The experiment is a toss of a two fair coins (or a double toss of a one coin). Then $\Omega = \{HH, HT, TH, TT\}$. Let X_3 be the indicator of different sides in the results, i.e. $X_3(HH) = X_3(TT) = 0$ and $X_3(HT) = X_3(TH) = 1$. Then $X_3 \sim \text{Bernoulli}(0.5)$;
- The experiment is a roll of a fair die. Then $\Omega = \{1, 2, 3, 4, 5, 6\}$. Now, let X_4 be 0, if the result will show odd number, and let it be 1, if the result is an even number. That is, $X_4(1) = X_4(3) = X_4(5) = 0$ and $X_4(2) = X_4(4) = X_4(6) = 1$. Then, clearly, $X_4 \sim \text{Bernoulli}(0.5)$;
- The experiment is again a roll of a fair die. Then $\Omega = \{1, 2, 3, 4, 5, 6\}$. Now, let X_5 be 0, if the result is ≤ 3 , and let it be 1, if the result is > 3 . That is, $X_5(1) = X_5(2) = X_5(3) = 0$ and $X_5(4) = X_5(5) = X_5(6) = 1$. Then, clearly, $X_5 \sim \text{Bernoulli}(0.5)$;
- The experiment is to pick a random number (uniformly) from $[0, 1]$. The probability measure is the length. Define $X_6(\omega) = 0$, if $\omega \in [0, 0.5)$ and $X_6(\omega) = 1$, if $\omega \in [0.5, 1]$. Then $X_6 \sim \text{Bernoulli}(0.5)$.

- The experiment is again to pick a random number (uniformly) from $[0, 1]$. The probability measure is the length. Define $X_7(\omega) = 0$, if $\omega \in [0, 0.2) \cup (0.7, 1]$ and $X_7(\omega) = 1$, if $\omega \in [0.2, 0.7]$. Then $X_7 \sim \text{Bernoulli}(0.5)$.

Another Supplement:

EXAMPLE, EXPLICITLY GIVEN R.V. AND ITS DISTRIBUTION: Assume X_n is given explicitly: let $\Omega = [0, 1]$, and the probability is given through $\mathbb{P}([a, b]) = b - a$ for any $[a, b] \subset \Omega$. Let X_n be the following r.v.:

$$X(\omega) = \begin{cases} 4, & \omega \in [0, 0.3] \\ -3, & \omega \in (0.3, 1]. \end{cases}$$

What can be said about the distribution of X ?

Clearly, X takes only 2 values, -3 and 4 , so X is discrete. The PMF of X will be:

$$\mathbb{P}(X = -3) = \mathbb{P}(\omega \in (0.3, 1]) = 1 - 0.3 = 0.7, \quad \text{and} \quad \mathbb{P}(X = 4) = \mathbb{P}(\omega \in [0, 0.3]) = 0.3 - 0 = 0.3.$$

So we can write the distribution of X in the following form:

Values of X	-3	4
$\mathbb{P}(X = x)$	0.7	0.3

So, in fact, the distribution of X is very familiar to us.

But think like this: if we will have the distribution of X in the table form above, we, unfortunately, cannot say anything about the particular values $X(\omega)$, we cannot reconstruct $X(\omega)$ (even we cannot know where is running ω , i.e., what is the Sample Space Ω).

Parametric inference and point estimation

One of the main problems in statistics is the following: we have a dataset x_1, \dots, x_n , a set of some observations, numbers (or a pair, tuple of numbers). We assume that this dataset comes as a realization of r.v.s X_1, \dots, X_n , coming from the same distribution F . Our aim is to get some information about F .

Why we are doing this? Because our aim, in general, is to get an information from the sample about the population. So our F describes our population. Unfortunately, usually we do not have the total population's distribution, or data collected from any individual from the population (census). So we think about sampling - to take a sample of some size n , and to collect information from n individuals.

EXAMPLE, INFO ABOUT POPULATION AND A SAMPLE: Assume we are interested in the probability of having a girl child, in Armenia. This probability is unknown, and nobody ever can know the exact value of that¹. What to do then? Maybe you have heard that the probability of having a girl child is a little bit larger than the probability of having a boy child (although any irén hargogh Armenian wants to have a boy child, ojakhi tsux@pahelu hamar). How we can estimate that probability? We can estimate that probability based on observations made in maternity hospitals or in zags-es. We can choose 1000 newborns and calculate the ratio of girls in that sample. And this will give some info about the probability we wanted to estimate.

If we will do the sampling, and collect information, we will have an n -tuple of numbers x_1, \dots, x_n . But we can think like this: this is just a random outcome of sampling - if we will do that sampling again, and collect information again from n individuals, maybe we will get another n numbers. That is, our observed sample is *one of the possible observations we could have*. So to model the situation, and take into account that we want to get an information about the population, about the process of generating that observed numbers, and not only about that observed particular sample, we can assume, before sampling, that the result of sampling is a sequence of r.v. X_1, X_2, \dots, X_n . In other words, to assess how good is our estimation, we need to consider the process of generation of the observations, not just one observation.

In short, we are modeling our dataset as a realization of a random sample X_1, \dots, X_n , i.e., as a realization (one of the possible realizations) of an n -tuple of IID r.v. X_1, \dots, X_n .

And, having our model, the problem of statistical inference is to recover the unknown distribution observing some finite number of samples from that distribution. So our general problem will be the following:

Problem: Assume that we have a sample X_1, X_2, \dots, X_n of IID random variables coming from the same distribution F :

$$X_1, X_2, \dots, X_n \stackrel{\text{IID}}{\sim} F.$$

The problem is to infer F .

Again, the idea is that we think about the data as having a random source, as coming from some random experiment (or repetitions of a random experiment), so we model it with random variables.

9.1 Random Samples and Statistics

So we start our story by defining a Random Sample:

Definition 9.1. *If we are given a sequence of IID r.v. X_1, X_2, \dots, X_n , then we will say that we have a **Random Sample** of size n .*

If we are given the values of X_1, X_2, \dots, X_n at some ω , $X_1(\omega) = x_1, \dots, X_n(\omega) = x_n$, then we say that x_1, \dots, x_n is a realization or an observation of our random sample.

What we want to do: as we have discussed above, one of the main problems in Statistics (as a Scientific direction) is to infer information about the distribution behind the data x_1, \dots, x_n . So, in general, we have only x_1, \dots, x_n , and we want to give some information. Of course, that information will depend on x_1, \dots, x_n only (if we do not have any other side info, say, if we do not have that our sample comes from an exponential distribution). So any numerical information we can get will be a function of x_1, \dots, x_n , say, of the form $g(x_1, \dots, x_n)$. Since our x_i -s are numbers (observations), then $g(x_1, \dots, x_n)$ will be a number too.

For example, we were calculating Sample Mean, Sample Variance etc. - they are all of the form $g(x_1, \dots, x_n)$, and all are numbers. Now, to obtain information about the general population, the distribution behind our data, we will use the Random Sample X_1, \dots, X_n , to obtain that information. So we will form functions of our random sample X_1, \dots, X_n , and work with them - this is called a **Statistics**²:

Definition 9.2. *Any function of a Random Sample X_1, \dots, X_n is called a **statistics**. So statistics is a r.v. of the form*

$$g(X_1, X_2, \dots, X_n).$$

In words, any information we can get from our random sample X_1, \dots, X_n is a statistics of the distribution behind our random sample.

Since $g(X_1, X_2, \dots, X_n)$ is a r.v., then we can talk about its distribution. The distribution of a statistic $g(X_1, X_2, \dots, X_n)$ is called the **Sampling Statistics**.

EXAMPLE, SAMPLING STATISTICS: Assume

$$X_1, \dots, X_n \stackrel{\text{IID}}{\sim} \text{Bernoulli}(p)$$

for some unknown $p \in [0, 1]$. The following r.v.'s are statistics:

$$A = X_1 + X_n, \quad \bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n}, \quad S_n = X_1 + \dots + X_n,$$

$$P_n = X_1 \cdot X_2 \cdot \dots \cdot X_n, \quad B = \sin(X_1 + X_2^2).$$

Say, if we want to find the sampling distribution of S_n , we can use the fact that the sum of n independent $\text{Bernoulli}(p)$ r.v.'s is Binomial with parameters (n, p) :

$$S_n = X_1 + \dots + X_n \sim \text{Binom}(n, p).$$

²In the third meaning of the word, as we have talked above.

Usually, it is not an easy and solvable-by-a-hand task to find the Sampling Distribution of a Statistics. We will not talk much about this (later we will consider the Sampling Distributions for some Statistics), but let me just add that one can use the Monte Carlo method to approximate that distribution by simulations.

9.2 Parametric and Nonparametric Models

When doing statistical inference, it is important to check first what kind of (prior) information we have about the unknown distribution F . Maybe we know (using some considerations) that our sample comes from the Exponential distribution, but we do not know the rate parameter for that distribution. In this case our problem will be to estimate that rate parameter. But maybe we do not have any additional information, and all we have that X_i -s are from the same distribution, so we will need to "recover" that distribution.

Usually, this kind of prior information comes from the nature of the problem we are solving. Say, if we want to find (in fact, estimate) the probability of having a girl child, of estimating the probability of producing defective detail in a factory, the probability of having a bad creditor, or the probability of voting for the candidate A, we will think about the Bernoulli model - we just have 2 possibilities, 0 or 1 (say, in the girl child problem we can assume that 1 corresponds to a girl, 0 to a boy), so Exponential model is not appropriate. Or, if we are modeling the number of calls during 1 hour in a taxi call center, or number of clicks on the webpage during some short period of time, the number of car accidents in a day, then nobody will think about the Bernoulli model, and the appropriate model will be the Poisson model. And so on - to use a good model, one needs to know which kind of phenomena can be modeled by which kind of probability distributions. We will not talk much about this, (un)fortunately. You can find in wiki pages for each distribution some examples where that distribution will be appropriate.

Usually one can divide the Statistics Subject into two main areas - Parametric and Non-Parametric Statistics (and also there is a Semi-Parametric one). The former one deals with parametric families of distributions, that is, here we assume that we know that our data comes from the fixed family of parametric distributions, and we need to estimate the unknown parameter. This area of statistics is widely studied, and we will talk mainly about it. The Non-Parametric Statistics deals with general distributions, not making any prior (or making just very general) assumptions.

9.3 Parametric Families of Distributions

In this part of our lectures, we are dealing with the Parametric Statistics, i.e., we will assume that our data comes from some parametric family of distributions.

Parametric family of distributions is a set of distributions of the form

$$\{\mathbb{P}_\theta : \theta \in \Theta\},$$

where θ is our parameter, Θ is the set of all possible parameters, \mathbb{P}_θ is the distribution with the parameter θ , given through its PD(M)³ $f(x|\theta)$ or CDF $F(x|\theta)$.

EXAMPLE, BERNOULLI FAMILY OF DISTRIBUTIONS: The Bernoulli Distributions family is given by

$$\{\text{Bernoulli}(p) : p \in [0, 1]\}$$

³We will write $f(x|\theta)$ to separate the **variable** x and the **parameter** θ .

Here our parameter is $\theta = p$, the set of parameters is $\Theta = [0, 1]$, and the distribution with the parameter p is Bernoulli(p).

If $X \sim \text{Bernoulli}(p)$, then the PMF of X is

$$f(x|p) = p^x \cdot (1 - p)^{1-x}, \quad x \in \{0, 1\}.$$

EXAMPLE, EXPONENTIAL FAMILY OF DISTRIBUTIONS: The Exponential Distributions family is given by

$$\{\text{Exp}(\lambda) : \lambda \in (0, +\infty)\}$$

Here our parameter is $\theta = \lambda$, the set of parameters is $\Theta = (0, +\infty)$, and the distribution with the parameter λ is $\text{Exp}(\lambda)$.

If $X \sim \text{Exp}(\lambda)$, then the PDF of X is

$$f(x|\lambda) = \begin{cases} \lambda \cdot e^{-\lambda x}, & x \geq 0 \\ 0, & x < 0 \end{cases}$$

EXAMPLE, NORMAL FAMILY OF DISTRIBUTIONS: The Normal Distributions family is a two-parameter family given by

$$\{\mathcal{N}(\mu, \sigma^2) : (\mu, \sigma^2) \in \mathbb{R} \times [0, +\infty)\}$$

Here our parameter is two-dimensional: $\theta = (\mu, \sigma^2)$, the set of parameters is $\Theta = \mathbb{R} \times [0, +\infty) \subset \mathbb{R}^2$, and the distribution with the parameter θ is $\mathcal{N}(\mu, \sigma^2)$.

If $X \sim \mathcal{N}(\mu, \sigma^2)$, then the PDF of X is

$$f(x|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot e^{-\frac{(x-\mu)^2}{\sigma^2}}, \quad x \in \mathbb{R}.$$

If we are considering a parametric family of distributions

$$\{\mathbb{P}_\theta : \theta \in \Theta\},$$

and we have $X \sim \mathbb{P}_\theta$, then the expected value, variance and other characteristics of X will depend on θ , in general. So we will denote them as $\mathbb{E}_\theta(X)$, $\text{Var}_\theta(X)$ etc.

EXAMPLE, PARAMETRIC FAMILY OF DISTRIBUTIONS:

If we consider the Bernoulli family $\{\text{Bernoulli}(p) : p \in [0, 1]\}$, and $X \sim \text{Bernoulli}(p)$, then $\mathbb{E}(X) = p$, so the expectation of X depends on the parameter p , and we write this as $\mathbb{E}_p(X) = p$, to underline the dependence on the parameter. Also, the variance of X will be

$$\text{Var}_p(X) = p(1 - p).$$

In general, if $X \sim \mathbb{P}_\theta$ and $f(x|\theta)$ is the PDF of \mathbb{P}_θ , then

$$\mathbb{E}_\theta(X) = \int_{-\infty}^{\infty} x \cdot f(x|\theta) dx, \quad \text{and} \quad \text{Var}_\theta(X) = \int_{-\infty}^{\infty} (x - \mathbb{E}_\theta(X))^2 \cdot f(x|\theta) dx.$$

9.4 Point Estimators and Estimates

Now, assume we have a parametric model (parametric family of distributions) $\{\mathbb{P}_\theta : \theta \in \Theta\}$. By this we mean that we have a family of distributions given by their CDFs or PDFs, depending on one or several parameters.

And assume we have a random sample

$$X_1, X_2, \dots, X_n \stackrel{\text{iid}}{\sim} \mathbb{P} \in \{\mathbb{P}_\theta : \theta \in \Theta\}$$

and our aim is to estimate \mathbb{P} . So we know that \mathbb{P} needs to be from that parametric model, and we know that all X_i -s have distribution \mathbb{P} , and we need to estimate \mathbb{P} . This means we need to estimate that $\theta \in \Theta$ for which $\mathbb{P} = \mathbb{P}_\theta$.

Let the set of parameters Θ be a subset of \mathbb{R}^p , so our parameter is p -dimensional (for simplicity, you can think of the case $p = 1$, when our parameters are one dimensional). Our point estimate for θ is a single vector from \mathbb{R}^p that is our guess for the actual value of θ .

Definition 9.3. Take any function⁴ with values in $\Theta \subset \mathbb{R}^p$:

$$g : \mathbb{R}^n \rightarrow \Theta.$$

The random vector

$$\hat{\theta} = \hat{\theta}_n = g(X_1, \dots, X_n)$$

is called an estimator for θ , and for any realisation x_1, x_2, \dots, x_n of X_1, \dots, X_n , the p -dim vector

$$g(x_1, \dots, x_n)$$

is called an estimate for θ .

So for any realization of X_1, \dots, X_n , we will get one estimate for θ , namely, $g(x_1, \dots, x_n)$. And if we will do sampling many times, all these estimates can be different.

REMARK, ESTIMATORS: It is important that the estimator cannot depend on the unknown parameter, simply because that parameter is unknown, and we want to estimate it, using the estimator, through the results of experiment, measurement, which will be known as soon as we will do the experiment, make observations.

Say, if we will take $g = X_1^3 + 2\theta X_2$, where θ is our parameter to estimate, then we will not be able to use the values of g to estimate θ . But, of course, we can use $g = X_1^3 + 2X_2$ as an estimator. And, if we will do the experiment and get observations for X_k -s, say, $X_1 = 2$ and $X_2 = 4$, then we will get the estimate $g = 2^3 + 2 \cdot 4$ for θ . Here I am not considering if our estimate is a good one or a bad one - we will talk about that soon.

REMARK, ESTIMATORS, ESTIMATES AND PARAMETERS: Please note that our unknown parameter θ is fixed, is a vector (or, in 1D, is just a real number), but $\hat{\theta}_n$ is a r.v., so estimator is a **random vector (variable)**. And, an estimate is again a number, and this time it is some fixed, known, number (obtained when plugging known observations into a known formula for the estimator).

To remember:

θ – our parameter, a constant, number (although unknown), not a r.v.

⁴Measurable function!

$$\hat{\theta} = \hat{\theta}_n = \hat{\theta}_n(X_1, X_2, \dots, X_n) - \text{estimator, a r.v.}$$

$$\hat{\theta} = \hat{\theta}_n = \hat{\theta}_n(x_1, x_2, \dots, x_n) - \text{estimate, a number, known}$$

9.5 Some Good (Desirable) Properties of Point Estimators

Estimators, defined the above way, as functions of random samples, are pretty much arbitrary chosen, and they can be very far from the actual values of θ . So we need to measure somehow how good is our estimator. For a "good" estimator, the estimates for all realisations $g(x_1, \dots, x_n)$ need to fall "close" to θ . That is, we want to have that the distribution of our estimator $g(X_1, \dots, X_n)$ be concentrated around θ .

Now we give some measures to calculate how good is our estimator. For simplicity, we assume $p = 1$, so $\Theta \subset \mathbb{R}$, our parameters are 1D.

Please note that at this point it will not be clear how or where from I am choosing/taking estimators. I will take different estimators and try to compare them, check if they possess some good properties etc. Later we will talk about some methods to generate good estimators.

9.5.1 Risk of the Estimator, Mean Squared Error

The first and very important measure for assessing how well our estimator is, which takes into account both the mean and the variance, is the Mean Squared Error or the Risk:

Definition 9.4. *The Mean Squared Error or the Quadratic Risk of the estimator $\hat{\theta}_n$ of θ is*

$$\text{MSE} = \text{MSE}(\hat{\theta}_n, \theta) = \text{Risk}(\hat{\theta}_n, \theta) = R(\hat{\theta}_n, \theta) = \mathbb{E}_\theta[(\hat{\theta}_n - \theta)^2].$$

In some sense, MSE is measuring how far can be the values of $\hat{\theta}_n$ from the real value of the parameter θ . Say, the best case will be when the MSE will be zero. But that can happen only if $\hat{\theta}_n \equiv \theta$, which is impossible, in general - we need to be able to calculate the values of estimators, but here we do not know the real, actual value of θ .

First of all, let us talk about comparing two estimators. Assume we have two estimators $\hat{\theta}_n^1$ and $\hat{\theta}_n^2$ for the same parameter θ .

Definition 9.5. *We say that the estimator $\hat{\theta}_n^1$ of θ is preferable to $\hat{\theta}_n^2$, another estimator of θ , if*

$$\text{MSE}(\hat{\theta}_n^1, \theta) \leq \text{MSE}(\hat{\theta}_n^2, \theta), \quad \forall \theta \in \Theta,$$

and there exists a θ s.t. $\text{MSE}(\hat{\theta}_n^1, \theta) < \text{MSE}(\hat{\theta}_n^2, \theta)$.

EXAMPLE, COMPARISON OF ESTIMATORS: Consider the following problem: we consider all distributions with finite means and with finite and fixed, known variance σ^2 , and our parameter is the mean:

$$X_1, X_2, \dots, X_n \stackrel{\text{iid}}{\sim} F,$$

and our parameter is the mean of F , $\mu = \mathbb{E}(X)$, where $X \sim F$. We want to estimate μ based on the random sample we have (or, more accurately, based on the observation from that random sample).

To that end, we can consider different estimators:

$$\hat{\mu}_n^1 = X_2, \quad \hat{\mu}_n^2 = \frac{X_1 + X_n}{2}, \quad \hat{\mu}_n^3 = \frac{X_1 + X_2 + X_3}{2}, \quad \hat{\mu}_n^4 = \frac{X_1 + X_2 + \dots + X_n}{n}.$$

Let's calculate risks for each estimator, and then choose the best one (to estimate μ) among these 4 estimators.

$$\begin{aligned} \text{Risk}(\hat{\mu}_n^1, \mu) &= \text{MSE}(\hat{\mu}_n^1, \mu) = \mathbb{E}_\mu((\hat{\mu}_n^1 - \mu)^2) = \mathbb{E}_\mu((X_2 - \mu)^2) \stackrel{\mathbb{E}(X^2) = \text{Var}(X) + [\mathbb{E}(X)]^2}{=} \text{Var}(X_2 - \mu) + \\ &(\mathbb{E}(X_2 - \mu))^2 = \text{Var}(X_2) + (\mathbb{E}(X_2) - \mu)^2 \stackrel{X_2 \sim F}{=} \text{Var}(X_2) + 0 = \sigma^2. \end{aligned}$$

....

Do not look at me (my notes) like this - I was warning you that my notes are incomplete ☺

Now, the idea of choosing and using a good estimator is that to choose an estimator with possible smallest Risk. Up to this point we have talked about how to compare two estimators, but we can also think about the minimal risk estimator. That is, one can try to find, among all estimators of θ , the one with the minimal risk, i.e.,

$$\hat{\theta}_n^* \in \underset{\hat{\theta}_n}{\text{argmin}} \text{MSE}(\hat{\theta}_n, \theta),$$

where $\hat{\theta}_n$ is running over all estimators of θ . Unfortunately, this problem is not solvable in the set of all estimators. (?! MP, add some words why is it unsolvable)

But, fortunately, there are some tools to work with risk and choose some good estimators in the sense of the risk.

Next section is introducing the notions of a bias and unbiasedness, and we will see the Bias-Variance decomposition of the risk there, and talk about choosing the Minimum Risk Unbiased Estimators.

REMARK, LOSS FUNCTION AND RISK: In general, one can measure how well is performing our estimator by using a loss function. One first defines the loss function

$$L : \Theta \times \Theta \rightarrow \mathbb{R},$$

with the properties

$$L(\theta_1, \theta_2) \geq 0 \quad \text{and} \quad L(\theta_1, \theta_2) = 0 \quad \text{iff} \quad \theta_1 = \theta_2,$$

then defines the average loss of the Risk of an estimator $\hat{\theta}_n$ to θ as

$$\text{Risk}(\hat{\theta}_n, \theta) = \text{Risk}_{\mathcal{L}}(\hat{\theta}_n, \theta) = \mathbb{E}_\theta[\mathcal{L}(\hat{\theta}_n, \theta)].$$

As examples of loss functions we can take

$$L(\theta_1, \theta_2) = |\theta_1 - \theta_2| \quad \text{or} \quad L(\theta_1, \theta_2) = (\theta_1 - \theta_2)^2.$$

The later loss function gives us the MSE, and the former one gives us the Risk

$$\text{Risk}(\hat{\theta}_n, \theta) = \mathbb{E}_\theta(|\hat{\theta}_n - \theta|).$$

9.5.2 Bias and Unbiasedness

Definition 9.6. The *bias* of estimator $\hat{\theta}_n$ of θ is

$$\text{bias}(\hat{\theta}_n, \theta) = \mathbb{E}_\theta(\hat{\theta}_n - \theta) = \mathbb{E}_\theta(\hat{\theta}_n) - \theta.$$

Here $\mathbb{E}_\theta(\hat{\theta}_n)$ is the expected value of $\hat{\theta}_n = \hat{\theta}_n(X_1, \dots, X_n)$, when we assume that the distribution behind every X_k is \mathbb{P}_θ .

Definition 9.7. We say that our estimator $\hat{\theta}_n$ (of the parameter θ) is **unbiased**, if

$$\text{bias}(\hat{\theta}_n, \theta) = 0 \quad \text{for any} \quad \theta \in \Theta. \quad (9.1)$$

If $\hat{\theta}_n$ is not unbiased, then we say that it is **biased**.

In other words, the estimator $\hat{\theta}_n$ is **unbiased**, if

$$\mathbb{E}_\theta(\hat{\theta}_n) = \theta \quad \text{for any} \quad \theta \in \Theta.$$

REMARK, UNBIASEDNESS OF THE ESTIMATOR: Here, in the definition of the unbiasedness, the equality of the bias to zero for **any value of** $\theta \in \Theta$ is very important. This is because actually we do not know the exact value of the parameter θ . Maybe our bias will be zero for some θ , and that θ is not the correct θ - that will give that for that incorrect θ , the estimator's values are concentrated around that θ , but this is not what we wanted. The property needs to hold for any $\theta \in \Theta$, and in that case we will be sure that it is true for our unknown parameter value θ .

In some sense, bias measures how accurate is our estimator. If our estimator is unbiased, then we have pretty accurate estimator, and if we will do a lot of samplings, and average the results, then we will be able to recover the true parameter with a high accuracy. On the other hand, if doing a lot of experiments is not possible or is very expensive (making observations is not so easy), then, having an unbiased estimator is not giving a confidence that the estimate will be close to the true one.

EXAMPLE, UNBIASEDNESS: For example, assume we want to estimate the mean lifetime of some type of light bulb. Let that mean lifetime be denoted by μ , so μ is our unknown parameter. Using the probabilistic model, we assume that the lifetime of a bulb (generic bulb of that type, not a particular one) is modeled by a r.v. X , and we want to estimate $\mathbb{E}(X) = \mu$. To that end, we can take some light bulbs of that type and record their lifetimes. Say, we will use the data for 5 light bulbs. Before recording the lifetimes, we will have a random sample X_1, X_2, X_3, X_4, X_5 , where X_k will show the lifetime of the k -th observed bulb. Until doing the actual observation, these are r.v.'s, with the same distribution (which is the distribution of X). Now, our aim was to estimate μ , using X_k , $k = 1, \dots, 5$. Here we need to take some estimator to estimate μ . Say, I am taking the estimator $\hat{\mu} = X_1$. This means that I will take the first observed lifetime as my estimate for μ . I know, this is not a good idea to estimate μ , but it is a good idea to explain unbiasedness 😊

If I will base my estimation only on one observation, that can give very incorrect results. Say, the first light bulb can go out of order in 1 week, just because that one was not a good one, or something happened with electricity or just by a chance.

Say, I am doing a sampling (choosing 5 light bulbs at random and recording their lifetimes), and the results are (in months):

$$x_1 = 2.3, x_2 = 12.1, x_3 = 9.4, x_4 = 11.1, x_5 = 14$$

Using my estimator, the estimate for μ will be

$$\hat{\mu} = x_1 = 2.3.$$

Well, not so perfect, supposedly.

Now, let us do a sampling again (again choosing 5 light bulbs at random and recording their lifetimes). Say, this time the results are:

$$x_1 = 9.1, x_2 = 10.0, x_3 = 14.2, x_4 = 16.1, x_5 = 12.7$$

This time my estimate will be

$$\hat{\mu} = x_1 = 9.1.$$

So, using 2 samples, I have obtained 2 estimates. And what the unbiasedness means? If I will continuously sampling of 5 lifetimes and calculate the corresponding estimates, and denote them by, say, $\hat{\mu}^1 = 2.3, \hat{\mu}^2 = 9.1, \hat{\mu}^3, \dots, \hat{\mu}^N, \dots$, then unbiasedness says that

$$\frac{\hat{\mu}^1 + \hat{\mu}^2 + \dots + \hat{\mu}^N}{N} \rightarrow \mu, \quad N \rightarrow +\infty.$$

So even if individual estimates (for an unbiased estimator) can give wrong and incorrect results, if we can do many samplings, calculate many estimates, in the average we will have the approximate value of the true mean μ !

R CODE, BIAS AND UNBIASEDNESS: In this example, we will assume the lifetime of the bulb follows a Gamma distribution.

```
g.shape <- 10
g.scale <- 1
x <- rgamma(5, shape = g.shape, scale = g.scale)
mu <- g.shape * g.scale
mu.hat <- x[1]
mu.hat
mu
#Supposedly, the estimate mu.hat is not too close to mu

# now, repeated sampling
for (i in 1:1000){
  x <- rgamma(5, shape = g.shape, scale = g.scale)
  mu.hat[i] <- x[1]
}

mean(mu.hat)
mu
#Here we need to have more close numbers
```

```
#Graphing the values of estimates
plot(mu.hat, type = "l", lwd = 1)
abline( h = mu, col = "red", lwd = 2)
#Boxplot, finally
boxplot(mu.hat, horizontal = T)
```

EXAMPLE, HOW MANY UNBIASED ESTIMATORS A PARAMETER CAN HAVE?: Assume we have 2 unbiased estimators for the parameter θ : $\hat{\theta}_n$ and $\tilde{\theta}_n$. In that case, for any $\alpha \in \mathbb{R}$, the estimator

$$\hat{\theta}_n^\alpha = \alpha \cdot \hat{\theta}_n + (1 - \alpha) \cdot \tilde{\theta}_n$$

will be an unbiased estimator. So we will have infinitely many unbiased estimators for θ .
The Proof is simple and is kindly left to the reader to enjoy the beauty of Math.

REMARK, BIAS AND UNBIASEDNESS: Assume the estimator $\hat{\theta}_n = g(X_1, \dots, X_n)$ is an unbiased estimator for the parameter θ .

When talking about doing a lot of experiments and averaging the results, mathematically we mean that if we will have **estimates**

$$\hat{\theta}_n^k = g(x_1^k, \dots, x_n^k), \quad k = 1, \dots, N$$

after making N samplings and obtaining samples x_1^k, \dots, x_n^k for the observation k , $k = 1, \dots, N$, then the average

$$\frac{\hat{\theta}_n^1 + \hat{\theta}_n^2 + \dots + \hat{\theta}_n^N}{N}$$

will be close to the true value, if N is large enough. This is, in fact, a consequence of the LLN, as, using the LLN we will obtain that for the sequence of **estimators**

$$\hat{\theta}_n^k = g(X_1^k, \dots, X_n^k), \quad k = 1, \dots, N$$

we will have

$$\frac{\hat{\theta}_n^1 + \hat{\theta}_n^2 + \dots + \hat{\theta}_n^N}{N} \rightarrow \mathbb{E}(g(X_1, \dots, X_n)) = \mathbb{E}_\theta(\hat{\theta}_n) = \theta$$

in the a.s. sense (SLLN) or in Probability (WLLN). So for large N , after plugging the observations, with high probability, we will obtain a value close to the limit, θ .

EXAMPLE, BIAS AND UNBIASED ESTIMATOR: Assume we have a random sample

$$X_1, X_2, \dots, X_n \sim \mathbb{P},$$

where \mathbb{P} is from some parametric family $\{\mathbb{P}_\theta : \theta \in \Theta\}$. We want to estimate the population mean $\mu = \mathbb{E}(X)$, where $X \sim \mathbb{P}$. We can take, say, the following estimator:

$$\hat{\mu}_1 = X_1,$$

and this will be an unbiased estimator for μ , since

$$\mathbb{E}(\hat{\mu}_1) = \mathbb{E}(X_1) = \mu.$$

But, maybe, this is not a good choice, since if we will have a single observation from X_1 , then the chance that it will be a good representative for μ , is small. And, of course, we are forgetting about other information we have, about X_2, \dots, X_n .

Another estimator is

$$\hat{\mu}_2 = \frac{X_1 + X_2 + X_3}{3},$$

and it is unbiased, since again $E(\hat{\mu}_2) = \mu$. Or

$$\hat{\mu}_3 = \frac{X_1 + X_2 + X_3 + \dots + X_n}{n} = \bar{X}_n$$

which is again an unbiased estimator, since $E(\hat{\mu}_3) = \mu$.

If we will take, say

$$\hat{\mu}_4 = \frac{X_1 + X_2 + X_3 + \dots + X_{n-1}}{n},$$

then we will get another estimator for μ . Using the SLLN, it is easy to see that for large n , $\hat{\mu}_4$ is close to μ , so if we will plug the realizations of X_k -s, then the result will be close to μ . But in this case our estimator will be biased, and the bias will be:

$$\text{bias}(\hat{\mu}_4, \mu) = \dots$$

Sorry, you need to calculate that by yourself!

So basically, we measure the concentration of our estimator $\hat{\theta}_n = g(X_1, \dots, X_n)$ by calculating the distance of its expected value from θ . For an unbiased estimator, in average, estimates are equal to θ .

R CODE, BIASEDNESS AND UNBIASEDNESS: !! Explain the code (Note to MP)

```
#Biasedness and unbiasedness of estimators

est1 <- c()
est2 <- c()
est3 <- c()
est4 <- c()
est5 <- c()
n <- 80
for (i in 1:5000){
  x <- rbinom(n, 1, prob = 0.3)
  est1[i] <- x[1]
  est2[i] <- mean(x[1:5])
  est3[i] <- sum(x[1:(n-4)])/n
  est4[i] <- mean(x)
  est5[i] <- (x[1]+x[2])/5
}

bins = seq(from = -0.1, to = 1.1, by = 0.02)
```

```

#dev.new()
windows(width = 7, height = 7, pointsize = 12, title = "Histograms")
par(mfrow = c(2,3))
hist(est1, breaks = bins)
abline(v = 0.3, col = "red", lwd = 2)
hist(est2, breaks = bins)
abline(v = 0.3, col = "red", lwd = 2)
hist(est3, breaks = bins)
abline(v = 0.3, col = "red", lwd = 2)
hist(est4, breaks = bins)
abline(v = 0.3, col = "red", lwd = 2)
hist(est5, breaks = bins)
abline(v = 0.3, col = "red", lwd = 2)

#dev.new()
windows(width = 7, height = 7, pointsize = 12, title = "Boxplots")
par(mfrow = c(2,3))
boxplot(est1, horizontal = T)
abline(v = 0.3, col = "red", lwd = 2)
boxplot(est2, horizontal = T)
abline(v = 0.3, col = "red", lwd = 2)
boxplot(est3, horizontal = T)
abline(v = 0.3, col = "red", lwd = 2)
boxplot(est4, horizontal = T)
abline(v = 0.3, col = "red", lwd = 2)
boxplot(est5, horizontal = T)
abline(v = 0.3, col = "red", lwd = 2)

mean(est1)
mean(est2)
mean(est3)
mean(est4)
mean(est5)

```

If we will increase n above (do that experiment, several times!), then we will see that the bias of the estimator 3 vanishes (approaches 0), so, if we are able to take a lot of observations, then, in fact, we can get approximately unbiased estimates. This property, called asymptotic unbiasedness, is defined below.

Definition 9.8. The estimator $\hat{\theta}_n$ is called *asymptotically unbiased* for θ , if

$$\text{bias}(\hat{\theta}_n, \theta) \rightarrow 0, \quad \text{for any } \theta \in \Theta.$$

One important Property:

Proposition 9.1. (Bias-Variance Decomposition) If $\hat{\theta}_n$ is an estimator for θ , then

$$\text{MSE}(\hat{\theta}_n, \theta) = \left(\text{bias}(\hat{\theta}_n, \theta) \right)^2 + \text{Var}_{\theta}(\hat{\theta}_n).$$

Corollary 9.1. If $\hat{\theta}_n^1$ and $\hat{\theta}_n^2$ are both unbiased estimators for unknown parameter θ , then $\hat{\theta}_n^1$ is preferable to $\hat{\theta}_n^2$ if and only if

$$\text{Var}(\hat{\theta}_n^1) \leq \text{Var}(\hat{\theta}_n^2), \quad \text{for all } \theta.$$

Corollary 9.2. Assume $\hat{\theta}_n$ is an unbiased estimator for θ , and assume

$$\text{Var}(\hat{\theta}_n) \rightarrow 0.$$

Then $\hat{\theta}_n$ is consistent.

Example: Assume $X_1, \dots, X_n \sim \text{Binom}(n, p)$, where $p \in (0, 1)$. Then

- $\bar{X} = \frac{X_1 + \dots + X_n}{n}$ is an unbiased estimator for parameter p ;
- There is no unbiased estimator for $\frac{1}{p}$.

REMARK, BIAS-VARIANCE DECOMPOSITION: Bias-Variance decomposition can be visualized like this ...

See Fig. 9.1




Fig. 9.1: Bias-Variance decomposition illustration, from Knuggette

REMARK, POINT ESTIMATION: Sometimes we are interested in estimation of not the parameter itself, but of some function of that parameter. Mathematically, we try to estimate $h(\theta)$, where θ is our parameter. In that case we can do the same thing as above - to introduce estimators for $h(\theta)$, say, $\widehat{h(\theta)} = g(X_1, \dots, X_n)$, and then to calculate the bias

$$\text{bias}(\widehat{h(\theta)}, h(\theta)) = \mathbb{E}_\theta(\widehat{h(\theta)}) - h(\theta).$$

Well, having that the average is exactly θ is good enough, but it is not showing how concentrated are our estimated around θ , since we can have a very big variance. So in order to improve the measure of how well our estimator is, we need to take into account also the variance.

Definition 9.9. *Standard Error of estimator $\hat{\theta}$ is*

$$SE(\hat{\theta}) = SD(\hat{\theta}) = \sqrt{\text{Var}_{\theta}(\hat{\theta})}$$

Usually, when statisticians give an estimate for the parameter, they also give the Standard Error - this is to show somehow how good is the estimate, in the sense that if the Standard Error is small enough, the variation in the values of the estimator is small, the variability of estimates will be small⁵. In fact, the Standard Error will depend on the unknown parameter (as the Variance $\text{Var}_{\theta}(\hat{\theta})$ depends on θ) - so statisticians report the Estimated Standard Error, which is the Standard Error, where the value of θ is substituted by its estimate.

EXAMPLE, STANDARD ERROR: Assume we are estimating the parameter p for the Bernoulli(p) parametric family based on the observation 0, 1, 1, 0, 0, 0, 1, 0.

We are modeling the situation by using instead of the observation a Random Sample X_1, \dots, X_8 , assuming that

$$X_1, \dots, X_8 \stackrel{\text{IID}}{\sim} \text{Bernoulli}(p),$$

and our observation is one of the possible realizations of the Random Sample. We use the estimator

$$\hat{p} = \frac{X_1 + X_2 + \dots + X_8}{8} = \bar{X}$$

(later we will see that this is a good estimator for p). Now, the estimate for p , using our observation, will be

$$\hat{p} = \frac{0 + 1 + 1 + 0 + 0 + 0 + 1 + 0}{8} = \frac{3}{8}.$$

Now we want to give the Standard Error for the estimator. It is easy to see that

$$\begin{aligned} \text{Var}_p(\hat{p}) &= \text{Var}_p\left(\frac{X_1 + X_2 + \dots + X_8}{8}\right) = \frac{1}{8^2} \cdot \text{Var}_p(X_1 + X_2 + \dots + X_8) \stackrel{X_k \text{ are Indep}}{=} \\ &= \frac{1}{8^2} \cdot (\text{Var}_p(X_1) + \text{Var}_p(X_2) + \dots + \text{Var}_p(X_p)) \stackrel{X_k \sim \text{Bernoulli}(p)}{=} \frac{p(1-p)}{8}, \end{aligned}$$

so

$$SE(\hat{p}) = \sqrt{\frac{p(1-p)}{8}}.$$

This depends on the unknown parameter, p , so it cannot be used straightforwardly. So we calculate the Estimated Standard Error, the value of SE at the estimate:

$$\widehat{SE}(\hat{p}) = \sqrt{\frac{\frac{3}{8}(1-\frac{3}{8})}{8}} = \frac{1}{8}\sqrt{\frac{15}{8}}.$$

So statisticians say: based on the observation we have, the estimate for the probability of success (i.e., 1) is $\frac{3}{8}$, with the Standard Error $\frac{1}{8}\sqrt{\frac{15}{8}}$. S Vas 1300\$.

⁵Well, of course, this is not saying that the estimate is close to the real value of the parameter, rather that using this estimator you will obtain values not far from each other, and any other estimate will be close to the obtained one. This is not saying that that values will be close to the actual, real value of the parameter! To ensure this, you need to use a good estimator!

9.5.3 Consistency

Unbiasedness is good, but not too good. Unfortunately, usually one does not know the exact distribution of $\hat{\theta}_n$, as one doesn't know the distribution \mathbb{P}_θ behind every X_k . So the calculation of the bias can be very hard or even impossible. And even if we will have an unbiased estimator, that can help us in the case when we can do resampling, calculate the values of the estimator many times, and then average the results (see above). Another bad news is that it is possible that no unbiased estimator will exist for a parameter.

So statisticians are not talking about unbiasedness only. They consider also other nice properties that estimators can have. Some properties considered are for large datasets, if one can have a large amount of data - and then one considers the *asymptotic* behaviour of estimators, the behaviour for large n , on limiting behaviour when $n \rightarrow +\infty$. And now we want to give the definition of the consistency, which is one of the good properties that estimator can possess, when the number of our data becomes very large.

Definition 9.10. A point estimator $\hat{\theta}_n$ of the parameter θ is called

- **consistent**, if $\hat{\theta}_n \xrightarrow{\mathbb{P}} \theta$ for any $\theta \in \Theta$;
- **strongly consistent**, if $\hat{\theta}_n \xrightarrow{\text{a.s.}} \theta$ for any $\theta \in \Theta$;
- **weakly or Mean Square consistent**, if $\hat{\theta}_n \xrightarrow{L^2} \theta$ for any $\theta \in \Theta$, i.e., if

$$\text{MSE}(\hat{\theta}_n, \theta) = \mathbb{E}_\theta((\hat{\theta}_n - \theta)^2) \rightarrow 0 \quad \forall \theta \in \Theta.$$

R CODE, CONSISTENCY OF ESTIMATORS: Here we are considering the parametric family $\{\text{Exp}(\lambda), \lambda > 0\}$. We assume

$$X_1, X_2, \dots, X_n \stackrel{\text{IID}}{\sim} \text{Exp}(\lambda)$$

for some $\lambda > 0$, and our aim is to estimate λ .

We will use 3 estimators:

$$\hat{\lambda}_n^1 = \frac{X_1 + X_2 + \dots + X_n}{n}, \quad \hat{\lambda}_n^2 = \frac{n}{X_1 + X_2 + \dots + X_n}, \quad \hat{\lambda}_n^3 = \frac{X_1 + X_3}{2}$$

It can be shown, that only $\hat{\lambda}_n^2$ is consistent. !!! Do this

```
#Consistency of an estimator
#We estimate the Exp(lambda) parameter lambda
#The estimator lambda2 = 1/mean(x) is consistent
```

```
realrate <-2
no.of.exp <-1000
sample.size <-4000
lambda1 <- c()
lambda2 <- c()
lambda3 <- c()
for (i in 1:no.of.exp){
  x <- rexp(sample.size, rate = realrate)
  lambda1[i] <- mean(x)
```

```

lambda2[i] <- 1/mean(x)
lambda3[i] <- (x[1]+x[3])/2
}

windows(7,4)
par(mfrow = c(2,2))
hist(lambda1)
hist(lambda2)
hist(lambda3)

eps <-0.1
#calculating P(|lambda_k - realrate|>eps)
err1 <- abs(lambda1-realrate)
prob1 <- sum(err1>eps)/no.of.exp

err2 <- abs(lambda2-realrate)
prob2 <- sum(err2>eps)/no.of.exp

err3 <- abs(lambda3-realrate)
prob3 <- sum(err3>eps)/no.of.exp

```

EXAMPLE, NORMAL MODEL WITH KNOWN VARIANCE: We consider a parametric model $\{\mathcal{N}(\mu, 2) : \mu \in \mathbb{R}\}$, and we want to estimate unknown parameter μ based on a random sample

$$X_1, X_2, \dots, X_n \stackrel{\text{IID}}{\sim} \mathcal{N}(\mu, 2).$$

We consider the following estimator:

$$\bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n}.$$

We have the following properties for our estimator:

Property 1: The estimator \bar{X}_n is unbiased⁶

$$\mathbb{E}_\mu(\bar{X}_n) = \mathbb{E}_\mu\left(\frac{X_1 + X_2 + \dots + X_n}{n}\right) = \frac{\mathbb{E}_\mu(X_1) + \mathbb{E}_\mu(X_2) + \dots + \mathbb{E}_\mu(X_n)}{n}.$$

Now, since $X_k \sim \mathcal{N}(\mu, 2)$, then⁷ $\mathbb{E}_\mu(X_k) = \mu$, so we will obtain

$$\mathbb{E}_\mu(\bar{X}_n) = \frac{n \cdot \mu}{n} = \mu,$$

for any $\mu \in \mathbb{R}$. This means that the estimator is unbiased.

Property 2: The estimator \bar{X}_n is strongly consistent: this easily follows from the Strong LLN:

$$\bar{X}_n \xrightarrow{\text{a.s.}} \mu, \quad n \rightarrow +\infty,$$

so it is also consistent (since strong convergence implies convergence in probability).

Property 3: Let's calculate the Quadratic risk of the estimator \bar{X}_n to our parameter⁸ μ :

$$\text{MSE}(\bar{X}_n, \mu) = \mathbb{E}_\mu[(\bar{X}_n - \mu)^2] = \mathbb{E}_\mu[(\bar{X}_n)^2] - 2\mu \cdot \mathbb{E}_\mu(\bar{X}_n) + \mu^2 \stackrel{\mathbb{E}_\mu(\bar{X}_n) = \mu}{=} \mathbb{E}_\mu[(\bar{X}_n)^2] - 2\mu^2 + \mu^2 = \mathbb{E}_\mu[(\bar{X}_n)^2] - \mu^2.$$

Now we use the equality⁹ $\mathbb{E}(X^2) = \text{Var}(X) + (\mathbb{E}(X))^2$, so

$$\begin{aligned} \mathbb{E}_\mu[(\bar{X}_n)^2] &= \text{Var}_\mu(\bar{X}_n) + (\mathbb{E}_\mu(\bar{X}_n))^2 = \text{Var}_\mu\left(\frac{X_1 + X_2 + \dots + X_n}{n}\right) + \mu^2 \stackrel{X_k \text{ are iid}}{=} \\ &\stackrel{X_k \text{ are iid}}{=} \frac{\text{Var}_\mu(X_1) + \text{Var}_\mu(X_2) + \dots + \text{Var}_\mu(X_n)}{n^2} + \mu^2 = \frac{n \cdot \text{Var}_\mu(X_1)}{n^2} + \mu^2 \stackrel{X_k \sim \mathcal{N}(\mu, 2)}{=} \frac{n \cdot 2}{n^2} + \mu^2 = \frac{2}{n} + \mu^2. \end{aligned}$$

So plugging this value into the formula above, we will get

$$\text{MSE}(\bar{X}_n, \mu) = \mathbb{E}_\mu[(\bar{X}_n)^2] - \mu^2 = \frac{2}{n}.$$

Clearly, $\text{MSE}(\bar{X}_n, \mu) \rightarrow 0$ as $n \rightarrow +\infty$. This also means that $\bar{X}_n \xrightarrow{L^2} \mu$, so our estimator \bar{X}_n is also weakly consistent. ■

EXAMPLE, NORMAL MODEL ONCE AGAIN, WITH KNOWN VARIANCE: We consider the same parametric model $\{\mathcal{N}(\mu, 2) : \mu \in \mathbb{R}\}$, and again we want to estimate unknown parameter μ based on a random sample

$$X_1, X_2, \dots, X_n \stackrel{\text{iid}}{\sim} \mathcal{N}(\mu, 2).$$

In this example we consider a little bit different estimator:

$$\tilde{X}_n = \frac{X_2 + X_3 + \dots + X_n}{n},$$

so we are skipping the first r.v. X_1 . In this case \tilde{X}_n is biased, but is strongly consistent, again because of the Strong LLN:

$$\tilde{X}_n = \frac{X_2 + X_3 + \dots + X_n}{n} = \frac{X_2 + X_3 + \dots + X_n}{n-1} \cdot \frac{n-1}{n}$$

and by the Strong LLN,

$$\frac{X_2 + X_3 + \dots + X_n}{n-1} \xrightarrow{\text{a.s.}} \mu \quad \text{and} \quad \frac{n-1}{n} \rightarrow 1,$$

so

$$\tilde{X}_n \xrightarrow{\text{a.s.}} \mu.$$

We recommend to check if \tilde{X}_n is weakly consistent and to calculate the quadratic risk of estimation, and also to see which of the estimates \bar{X}_n from the previous example and \tilde{X}_n is preferable. ■

EXAMPLE, NORMAL MODEL ONCE AGAIN, WITH KNOWN MEAN: We consider now the model $\{\mathcal{N}(0, \theta) : \theta \in [0, +\infty)\}$, and we want to estimate the unknown variance θ having a random sample

$$X_1, X_2, \dots, X_n \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \theta).$$

We consider three estimators:

$$\hat{\theta}_n^1 = \frac{\sum_{k=1}^n (X_k - 0)^2}{n}, \quad \hat{\theta}_n^2 = \frac{\sum_{k=1}^n (X_k - \bar{X})^2}{n} \quad \text{and} \quad \hat{\theta}_n^3 = \frac{\sum_{k=1}^n (X_k - \bar{X})^2}{n-1}.$$

Here

$$\bar{X} = \bar{X}_n = \frac{X_1 + X_2 + \dots + X_n}{n}.$$

...

... The rest of consideration you need to do by yourself, so you need to study the properties of these estimators by yourself. Nice exercise before the Midterm.

REMARK, UNBIASEDNESS AND CONSISTENCY: We have talked above about two nice properties of estimators - Unbiasedness and Consistency. We have interpreted Unbiased Estimator as an estimator which gives accurate results in the sense that if we will do a lot of resamplings, calculate the values of the estimator and average them, then we will get a close approximation to the real value of the parameter. And consistent estimator is working well on large datasets. So what is the relation between these two notions - if we can do resamplings and results averaging, isn't it the same as to take a large dataset once?

In fact, no. Say, we can do a lot of resamplings even when our dataset is not big enough. Another point (or, maybe, the same one) is that taking a large sample means we will take each individual just once. But if we are doing resamplings, we can have the same individual in different samples. Say, we want to calculate the mean weight of the Yerevan's populations. To that end, we choose at random 1000 persons in Yerevan, ask about their weights, average them, then, on the next day, we choose again 1000 persons in Yerevan, get their weights, average them etc. But it can happen that the same person will be in both cases.

9.5.4 Choosing Good Estimators, Efficiency

Above we have defined several desirable properties of estimators. The idea is that before doing the actual estimation, one needs to be sure that he/she is using a "nice" estimator, which will give reliable, trustworthy results. Let us summarize the properties we have considered and their choices:

Assume we are estimating a parameter based on some dataset of observations. Then we consider several cases

A. If we have a large dataset, then:

- A1. If we can do resamplings, then we can use an asymptotically unbiased estimator: to estimate our parameter, we can calculate estimates for many samples and average the obtained values - that average will supposedly be close to the real value of the parameter;
- A2. If resampling is not available or preferable, then we can use a consistent estimator: to estimate a parameter, we can just calculate the estimate for one, large sample - the obtained value will be close to the real value of the parameter with high probability;

B. If we do not have a large dataset, then:

- B1. If we can do resamplings, then we can use an unbiased estimator: as above, to estimate our parameter, we can calculate estimates for many samples and average the obtained values - that average will supposedly be close to the real value of the parameter;

- B2. If resampling is not available or preferable, then we can use an estimator with a small Risk: to estimate a parameter, we can just calculate the estimate for one sample - the obtained value will be close to the real value of the parameter, if the risk is small.

Of course, an estimator can share all the above good properties, or, at least, some of them, and that will be a great estimator.

Soon we will talk about how to find/get estimators with good properties, we will consider Maximum Likelihood Method, Method of Moments and Bayesian Methods to obtain good estimators. By now, we know how to prove that the given estimator is unbiased, asymptotically unbiased or consistent. The only unexplained yet part is how to show that an estimator has a small Risk. And now we are going to talk about that.

As we have stated above, it is not always possible to find an estimator with the minimal Risk. So we restrict our attention to some subclass of estimators - we will consider only unbiased estimators, and try to find the one with the minimal risk.

Now assume $\hat{\theta}_n$ is an unbiased estimator for θ . Then, by the Bias-Variance Decomposition,

$$\text{Risk}(\hat{\theta}_n, \theta) = \text{MSE}(\hat{\theta}_n, \theta) = \text{Var}(\hat{\theta}_n).$$

So the problem of finding the estimator with the minimal Risk, for unbiased estimators, is equivalent to the problem of finding the estimator with the minimal Variance. And we have a new definition for that:

Definition 9.11. We say that the estimator $\hat{\theta}_n$ is **efficient** or a **MVUE** (Minimum Variance Unbiased Estimator) for the parameter θ , if¹⁰

- $\hat{\theta}_n$ is an unbiased estimator for θ ;
- $\hat{\theta}_n$ has the minimal variance among all unbiased estimators of θ , i.e., for any unbiased estimator $\tilde{\theta}_n$ of θ ,

$$\text{Var}_{\theta}(\hat{\theta}_n) \leq \text{Var}_{\theta}(\tilde{\theta}_n), \quad \forall \theta \in \Theta.$$

Having an efficient estimator is a nice thing everybody will dream for. So let us talk a little bit how to find efficient estimators.

Unfortunately, we cannot check the efficiency of an estimator by the definition, because we need to compare the variance of the estimator with variances of all unbiased estimators. Fortunately, we have a partial method that is giving a sufficient condition for efficiency, the Cramer-Rao theorem. Before stating the theorem, we need to give the idea of the Fisher information, which we will use in the Cramer-Rao's theorem statement.

9.5.5 Fisher Information

Assume we have a parametric family of distributions \mathbb{P}_{θ} , $\theta \in \Theta$, and $f(x|\theta)$ is the PD(M)F of \mathbb{P}_{θ} .

Definition 9.12. The following quantity is called **the Fisher Information** of the parametric family \mathbb{P}_{θ} :

$$I(\theta) = -\mathbb{E} \left(\frac{\partial^2}{\partial \theta^2} \ln f(X|\theta) \right) = \mathbb{E} \left[\left(\frac{\partial}{\partial \theta} \ln f(X|\theta) \right)^2 \right].$$

So to calculate the Fisher Information, you need to do the following steps:

¹⁰In some textbooks, efficiency is defined in other way, as an unbiased estimator which attains the Cramer-Rao lower bound, see the theorem above.

- Method 1:**
- Plug the r.v. X into the PD(M)F, so you will have $f(X|\theta)$;
 - Calculate the Natural Logarithm of the obtained r.v., $\ln f(X|\theta)$;
 - Calculate the second order partial derivative of the logarithm w.r.t θ , $\frac{\partial^2}{\partial \theta^2} \ln f(X|\theta)$;
 - Calculate the expected value of the obtained r.v.;
 - DO NOT FORGET to put a "-" sign in front of the obtained quantity.

- Method 2:**
- Plug the r.v. X into the PD(M)F, so you will have $f(X|\theta)$;
 - Calculate the Natural Logarithm of the obtained r.v., $\ln f(X|\theta)$;
 - Calculate the first order partial derivative of the logarithm w.r.t θ , $\frac{\partial}{\partial \theta} \ln f(X|\theta)$;
 - Calculate the square of the obtained r.v., $\left(\frac{\partial}{\partial \theta} \ln f(X|\theta) \right)^2$;
 - Calculate the expected value of the obtained r.v..

EXAMPLE, FISHER INFORMATION FOR THE Bernoulli(p) FAMILY: Tra-la-la
Explanation! Why information?

EXAMPLE, FISHER INFORMATION FOR THE $\mathcal{N}(\mu, \sigma^2)$ FAMILY: Oh-la-la
Explanation! Why information? Graphs!

EXAMPLE, FISHER INFORMATION FOR THE Exp(λ) FAMILY: Tra-la-la
Explanation! Why information?

REMARK, ANOTHER EXPRESSION FOR THE FISHER INFORMATION: It is easy to see that (under the regularity conditions)

$$\mathbb{E} \left(\frac{\partial}{\partial \theta} \ln f(X|\theta) \right) = 0.$$

This is because we can use the differentiation under the integral sign formula (assuming X is continuous) :

$$\begin{aligned} \mathbb{E} \left(\frac{\partial}{\partial \theta} \ln f(X|\theta) \right) &= \int_{-\infty}^{+\infty} \frac{\partial}{\partial \theta} \ln f(x|\theta) \cdot f(x|\theta) dx \stackrel{(\ln f)' = \frac{f'}{f}}{=} \int_{-\infty}^{+\infty} \frac{\frac{\partial}{\partial \theta} f(x|\theta)}{f(x|\theta)} \cdot f(x|\theta) dx = \\ &= \int_{-\infty}^{+\infty} \frac{\partial}{\partial \theta} f(x|\theta) dx = \frac{\partial}{\partial \theta} \int_{-\infty}^{+\infty} f(x|\theta) dx \stackrel{\int f=1}{=} \frac{\partial}{\partial \theta} (1) = 0. \end{aligned}$$

This means that the Fisher information is equal to:

$$I(\theta) = \mathbb{E} \left[\left(\frac{\partial}{\partial \theta} \ln f(X|\theta) \right)^2 \right] = \text{Var} \left(\frac{\partial}{\partial \theta} \ln f(X|\theta) \right).$$

Paho!!! Btw, the function

$$\frac{\partial}{\partial \theta} \ln f(x|\theta)$$

is called the score function.

9.5.6 Cramer-Rao Lower Bound (Cramer-Rao Inequality)

Now, we have all necessary tools to give the celebrated¹¹ Cramer-Rao theorem:

Theorem 9.1. (Cramer-Rao) Assume we have a Random Sample

$$X_1, X_2, \dots, X_n \stackrel{\text{IID}}{\sim} \mathbb{P}_\theta$$

and the Fisher Information for the family \mathbb{P}_θ is $I(\theta)$. Assume also that $\hat{\theta}_n$ is an unbiased estimator for θ obtained from our Random Sample. Then, under some regularity conditions¹² on the family \mathbb{P}_θ ,

$$\text{Var}_\theta(\hat{\theta}_n) \geq \frac{1}{n \cdot I(\theta)}.$$

What this theorem gives us - this gives a lower bound on the variance of *every* unbiased estimator. We know from the Bias-Variance decomposition, that

$$\text{MSE}(\hat{\theta}_n, \theta) = \left(\text{bias}(\hat{\theta}_n, \theta) \right)^2 + \text{Var}_\theta(\hat{\theta}_n).$$

We want to make MSE as small as possible, to choose the estimator $\hat{\theta}_n$ with the minimal MSE. The best will be to have an unbiased estimator (so $\text{bias}(\hat{\theta}_n) = 0$), and an estimator with a *veeery* small Variance $\text{Var}(\hat{\theta}_n)$. But Cramer and Rao, joining their voices, say that you cannot find an unbiased estimator with a *veeeery* small Variance - we have a fundamental and unbreakable lower bound for any unbiased estimator's variance - that variance cannot be smaller than $\frac{1}{nI(\theta)}$. What a pity! Alas, life is not so simple...

But, importantly, this theorem gives us a sufficient condition to have an efficient estimator:

Corollary 9.3. If $\hat{\theta}_n$ is an unbiased estimator and

$$\text{Var}_\theta(\hat{\theta}_n) = \frac{1}{n \cdot I(\theta)},$$

then $\hat{\theta}_n$ is efficient.

EXAMPLE, EFFICIENT ESTIMATOR FOR p IN Bernoulli(p) MODEL: Here we need to have that example from one of our lecture

EXAMPLE, EFFICIENT ESTIMATOR FOR μ IN $\mathcal{N}(\mu, \sigma^2)$ MODEL: Here we need to have that example from one of our lecture

REMARK, ABOUT THE EFFICIENCY:

Yesterday,
all my troubles seemed so far away.
Now it looks as though they're here to stay.
Oh I believe in yesterday.

My song for this Saturday. Will have a lot of tests to grade.
Thank you for your attention! Zanares.

¹¹And fundamental, and groundbreaking, and incredible, and hiasqanch, and chnashkharhik, and ...

¹²See ...